

Do Frontier LLMs Truly Understand Smart Contract Vulnerabilities?

Anonymous ACL submission

Abstract

Frontier large language models achieve remarkable performance on code understanding tasks (Claude Opus 4.5: 74.4% on SWE-bench, Gemini Pro Preview: 74.2%), yet their capacity for smart contract security remains unclear. Can they genuinely reason about vulnerabilities, or merely pattern-match against memorized exploits? We introduce **BlockBench**, a benchmark designed to answer this question, revealing heterogeneous capabilities. While some models demonstrate robust semantic understanding, most exhibit substantial surface pattern dependence.

1 Introduction

Smart contract vulnerabilities represent one of the most costly security challenges in modern computing. As shown in Figure 1, cryptocurrency theft has resulted in over \$14 billion in losses since 2020, with 2025 already reaching \$3.4 billion, the highest since the 2022 peak (Chainalysis, 2025). The Bybit breach alone accounted for \$1.5 billion, while the Cetus protocol lost \$223 million in minutes due to a single overflow vulnerability (Yellow Research, 2025).




Figure 1: Annual cryptocurrency theft losses (2020–2025). Data from Chainalysis.

Meanwhile, large language models have achieved remarkable success on programming tasks. Frontier models now pass technical interviews, generate production code, and identify bugs

across diverse codebases. This raises a natural question: *can these models apply similar expertise to blockchain security?* And if they can, *are they genuinely reasoning about vulnerabilities, or merely pattern-matching against memorized examples?*

This distinction matters. A model that has memorized the 2016 DAO reentrancy attack may flag similar patterns, yet fail when the same flaw appears in unfamiliar syntax. We introduce **BlockBench**, a benchmark designed to answer this question. Our contributions include:

1. **BlockBench**, comprising 263 Solidity vulnerability samples with systematic contamination control and gold standard examples from recent professional security audits.
2. **Composite evaluation metrics** distinguishing genuine understanding from memorization, validated through multi-configuration sensitivity analysis (Spearman’s $\rho=1.000$).
3. **Systematic assessment** revealing 58% best-case detection on mixed samples collapsing to 20% on uncontaminated professional audits, exposing heterogeneous robustness and accuracy-understanding gaps across models.

2 Related Work

Traditional Smart Contract Analysis. Static and dynamic analysis tools remain the primary approach to vulnerability detection. Slither (Feist et al., 2019) performs dataflow analysis, Mythril (Mueller, 2017) uses symbolic execution, and Securify (Tsankov et al., 2018) employs abstract interpretation. These tools achieve reasonable precision on well-defined vulnerability classes but exhibit significant false positive rates and limited coverage of complex semantic flaws (Durieux et al., 2020).

LLM-Based Vulnerability Detection. Recent work explores LLMs for smart contract analysis. GPTLens (Hu et al., 2023) employs adversarial auditor-critic interactions, while PropertyGPT (Liu

et al., 2024) combines retrieval-augmented generation with formal verification. Fine-tuned models achieve over 90% accuracy on benchmarks (Hosain et al., 2025), though performance degrades substantially on real-world contracts (Ince et al., 2025).

Benchmark Datasets. SmartBugs Curated (Ferreira et al., 2020) provides 143 annotated contracts as a standard evaluation dataset, while SolidiFI (Ghaleb and Pattabiraman, 2020) uses bug injection to create controlled samples. Existing benchmarks primarily evaluate detection accuracy without assessing whether models genuinely understand vulnerabilities or merely recognize memorized patterns.

LLM Robustness and Memorization. Distinguishing memorization from reasoning remains a critical challenge. Models exhibit high sensitivity to input modifications, with performance drops of up to 57% on paraphrased questions (Sánchez Salido et al., 2025). Wu et al. (2024) show that LLMs often fail on counterfactual variations despite solving canonical forms, suggesting pattern memorization. Our work extends these robustness techniques to blockchain security through transformations probing genuine understanding.

3 BlockBench

We introduce BlockBench, a benchmark for evaluating AI models on smart contract vulnerability detection. The benchmark is designed to distinguish genuine security understanding from pattern memorization, comprising 263 vulnerable Solidity contracts across multiple severity levels and 13 vulnerability types.

Let \mathcal{D} represent the dataset, where $\mathcal{D} = \{(c_i, v_i, m_i)\}_{i=1}^{263}$. Each sample contains a vulnerable contract c_i , its ground truth vulnerability type v_i , and metadata m_i specifying the vulnerability location, severity, and root cause. We partition \mathcal{D} into three disjoint subsets, $\mathcal{D} = \mathcal{D}_{\text{DS}} \cup \mathcal{D}_{\text{TC}} \cup \mathcal{D}_{\text{GS}}$, each targeting a distinct evaluation objective (Table 1).

Table 1: BlockBench composition spanning Critical, High, Medium, and Low severity.

Subset	N	Sources
Difficulty Stratified	179	SmartBugs, ToB
Temporal Contam.	50	DeFiHackLabs
Gold Standard	34	Spearbit, C4

Difficulty Stratified. \mathcal{D}_{DS} draws from established vulnerability repositories including SmartBugs Curated (Ferreira et al., 2020), Trail of Bits’ Not So Smart Contracts (Trail of Bits, 2018), and DeFiVulnLabs (SunWeb3Sec, 2023). Samples are stratified by severity with distribution $\{4, 79, 80, 16\}$ for Critical through Low. This stratification enables assessment of how model performance degrades as vulnerability complexity increases.

Temporal Contamination. \mathcal{D}_{TC} reconstructs well-known exploits from DeFiHackLabs (SunWeb3Sec, 2024) and the REKT Database (REKT Database, 2023), including Nomad Bridge (\$190M), Beanstalk (\$182M), and Curve Vyper (\$70M). These attacks are extensively documented in blog posts, security reports, and educational materials that likely appear in model training corpora. High performance on \mathcal{D}_{TC} may therefore reflect memorization of attack patterns rather than genuine vulnerability understanding.

Gold Standard. \mathcal{D}_{GS} derives from professional security audits by Spearbit (Spearbit, 2025), MixBytes (MixBytes, 2025), and Code4rena (Code4rena, 2025) conducted after September 2025. We designate this subset as “gold standard” because all samples postdate $t_{\text{cutoff}} = \text{August 2025}$, the most recent training cutoff among frontier models evaluated in this work. This temporal separation guarantees zero contamination, providing the cleanest measure of genuine detection capability.

Coverage. BlockBench spans 13 vulnerability classes. Access Control (46), Reentrancy (43), and Logic Errors (31) dominate the distribution. \mathcal{D}_{TC} emphasizes oracle manipulation and access control. \mathcal{D}_{GS} focuses on subtle logic errors. \mathcal{D}_{DS} provides broad coverage across classical patterns.

4 Methodology

Our evaluation methodology comprises four phases: adversarial transformation, model evaluation, automated judgment, and metrics computation. Figure 2 illustrates the complete pipeline.




Figure 2: BlockBench evaluation pipeline.

4.1 Adversarial Transformations

To distinguish memorization from understanding, we apply semantic-preserving transformations that systematically remove surface cues while preserving vulnerability semantics. For each contract $c \in \mathcal{D}$, we generate variants $\{\mathcal{T}_k(c)\}$ satisfying $\mathcal{V}(\mathcal{T}(c)) = \mathcal{V}(c)$, where \mathcal{V} extracts vulnerability semantics.

Sanitization (sn) removes security hints from identifiers and comments through 280+ pattern replacements while maintaining natural code style. **No-Comments (nc)** strips all documentation. **Chameleon (ch)** replaces blockchain terminology with domain-shifted vocabulary (medical, gaming themes). **Shapeshifter (ss)** applies multi-level obfuscation from identifier renaming (L2) to control flow obscuration (L3). This pipeline generates 1,343 variants from 263 base samples. Complete transformation specifications appear in Appendix B.

4.2 Evaluation Protocol

We evaluate six frontier models (Claude Opus 4.5, GPT-5.2, Gemini 3 Pro, Grok 4, DeepSeek v3.2, Llama 3.1 405B) using three prompt types. *Direct* requests structured JSON analysis. *Naturalistic* provides informal review requests. *Adversarial* includes misleading context claiming prior audit approval. All models use consistent parameters (temperature 0, max tokens 8192). Prompt templates appear in Appendix C.

4.3 Automated Judgment

Mistral Medium 3 serves as LLM judge, evaluating responses against ground truth. The judge classifies findings as TARGET_MATCH, BONUS_VALID, or invalid (HALLUCINATED, MISCHARACTERIZED, SECURITY_THEATER). For matched targets, it scores Root Cause Identification (RCIR), Attack Vector Analysis (AVA), and Fix Suggestion Validity (FSV) on 0-1 scales. Human evaluation of 20 responses validates reliability ($\kappa=0.91$ verdict agreement, $\rho=0.87$ correlation). Complete judge protocol and classification criteria appear in Appendix D.

4.4 Metrics

We rank models by *Target Detection Rate* (TDR), the proportion of samples where the documented vulnerability was correctly identified with both type and location accuracy. *Lucky Guess Rate* measures

correct verdicts without target identification. *Finding Precision* computes the proportion of reported findings that are correct. *Reasoning Quality* averages RCIR, AVA, and FSV scores for successfully identified targets.

We report *Security Understanding Index* (SUI) as a weighted composite: $SUI = 0.40 \cdot TDR + 0.30 \cdot \text{Reasoning} + 0.30 \cdot \text{Precision}$. Sensitivity analysis across five weight configurations confirms perfect ranking stability (Spearman's $\rho=1.000$). Complete metric definitions and sensitivity analysis appear in Appendix F and E.

5 Results

We evaluate six frontier models on 58 Solidity vulnerability samples across Temporal Contamination (TC), Gold Standard (GS), and Difficulty Stratified (DS) subsets.

5.1 Overall Performance




Figure 3: Target Detection Rate across all models. Best performer achieves 58% detection, while highest accuracy (88%) corresponds to lowest TDR (18%).

Table 2 and Figure 3 present aggregate performance ranked by Target Detection Rate (TDR). Gemini 3 Pro achieves highest detection (58%), followed by GPT-5.2 (56%) and Claude Opus 4.5 (53%). When combining detection, reasoning quality, and finding precision into the Security Understanding Index (SUI), GPT-5.2 ranks first (0.746) due to superior finding precision (77%), followed by Gemini 3 Pro (0.734).

Llama 3.1 405B exhibits the most severe accuracy-TDR gap: 88% accuracy yet only 18% TDR, correctly classifying vulnerable samples without identifying specific vulnerability types or locations. This 70-percentage-point discrepancy demonstrates that binary classification metrics inadequately measure security understanding.

Table 2: Overall performance ranked by Target Detection Rate. Best values bold.

Model	TDR	SUI	Acc	RCIR	AVA	FSV	Findings
Gemini 3 Pro	57.6	0.734	93.9	0.97	0.97	0.95	2.6
GPT-5.2	55.9	0.746	75.0	0.97	0.98	0.97	2.4
Claude Opus 4.5	52.9	0.703	83.8	0.98	0.99	0.97	3.5
Grok 4	44.1	0.677	69.1	0.98	1.00	0.97	2.1
DeepSeek v3.2	38.2	0.599	82.4	0.91	0.92	0.86	3.0
Llama 3.1 405B	17.9	0.393	88.1	0.88	0.90	0.83	2.0

All models achieving target detection show strong reasoning quality ($RCIR/AVA/FSV \geq 0.95$), with minimal variation in explanation quality across top performers.

5.2 Gold Standard Performance

Gold Standard samples from post-September 2025 audits guarantee zero temporal contamination. Performance drops substantially: Claude Opus 4.5 leads with 20% TDR, followed by Gemini 3 Pro (11%), GPT-5.2 (10%), and Grok 4 (10%). DeepSeek v3.2 and Llama 3.1 405B detect zero targets. All models experience 34-50 percentage point drops from overall to Gold Standard performance.

5.3 Transformation Robustness

Sanitization. Neutralizing security-suggestive identifiers causes variable degradation. On temporal contamination samples, top models achieve 60% TDR on baseline versions. Sanitization impacts differ: GPT-5.2 and DeepSeek v3.2 maintain performance, while Grok 4 drops 40pp, exposing varying reliance on lexical cues.

Domain Shift. Replacing blockchain terminology with medical vocabulary shows mixed impact. Performance ranges 20-60% TDR across models, with GPT-5.2 maintaining 60% detection while others show 20-50% degradation.

Prompt Framing. Performance varies significantly across direct, adversarial (claiming prior audit approval), and naturalistic prompts. Gemini 3 Pro and GPT-5.2 demonstrate robustness with 18-21pp drops from direct to non-direct prompts. Claude Opus 4.5 and DeepSeek v3.2 show larger degradation (21-39pp), while Llama 3.1 405B exhibits inconsistent behavior (adversarial: 0%, naturalistic: 25%).

5.4 Human Validation

Two security experts independently reviewed 20 responses. Inter-rater agreement: verdict $\kappa=0.91$, type match $\kappa=0.84$, reasoning $\kappa=0.78$. Human-

judge correlation: $\rho=0.87$ ($p<0.001$), 85% agreement, validating automated evaluation.

6 Discussion




Figure 4: Security Understanding Index trajectory across progressive transformations of TC samples. GPT-5.2 maintains near-constant performance while most models degrade, revealing varying degrees of surface pattern reliance.

Understanding versus Memorization. Figure 4 reveals heterogeneous robustness across models. GPT-5.2 maintains stable SUI ($78.0 \rightarrow 77.2$) through sanitization, domain shifts, and obfuscation, demonstrating genuine semantic understanding. In contrast, DeepSeek v3.2 degrades 19.7 points ($78.7 \rightarrow 59.0$), indicating surface pattern dependence. Most models exhibit intermediate behavior, leveraging lexical cues when available while retaining partial structural understanding (Chen et al., 2021; Wu et al., 2024). This heterogeneity suggests current training methods produce inconsistent abstraction capabilities across architectures (Sánchez Salido et al., 2025). While genuine security understanding is demonstrably possible, most frontier models have not achieved it.

Measurement Inadequacy. The accuracy-TDR

gap exposes fundamental metric limitations. Llama 3.1 405B achieves 88% accuracy yet only 18% TDR, correctly classifying samples as vulnerable without identifying specific flaw types or locations (Jimenez et al., 2024). For security practitioners requiring actionable findings, binary classification provides insufficient value. Effective evaluation must measure precise vulnerability localization, not merely anomaly detection.

Practical Implications. Current frontier models cannot serve as autonomous auditors. Best performance reaches 58% detection with substantial Gold Standard degradation (20% maximum). However, complementary strengths suggest ensemble potential: Grok 4 offers breadth, GPT-5.2 provides consistency, Claude delivers explanation quality. Workflows positioning LLMs as assistive tools with mandatory expert review align capabilities with current limitations (Hu et al., 2023; Ince et al., 2025).

Ethical Considerations. BlockBench poses dual-use risks: adversarial prompts demonstrate methods that could suppress detection, while detailed vulnerability documentation may assist malicious actors. We justify public release on several grounds: adversarial robustness represents a fundamental requirement for security tools, malicious actors will discover these vulnerabilities regardless, and responsible disclosure enables proactive mitigation. All samples derive from already-disclosed vulnerabilities and public security audits, ensuring no novel exploit information is revealed. Practitioners should avoid over-reliance on imperfect tools, as false negatives create security gaps while false confidence may reduce manual review rigor.

Limitations and Future Work. Our evaluation uses 263 samples with 10 Gold Standard examples. We assess zero-shot prompting exclusively. Chain-of-thought reasoning or retrieval augmentation may improve performance. Future work should expand sample diversity across blockchain ecosystems, develop sanitization-resistant analysis using control flow graphs, and explore hybrid LLM-verification architectures (Liu et al., 2024).

7 Conclusion

BlockBench evaluates whether frontier LLMs genuinely understand smart contract vulnerabilities or merely pattern-match against training data. Our assessment of six models on 263 Solidity samples reveals substantial limitations. Best performance reaches 58% detection on mixed samples, collaps-

ing to 20% on uncontaminated Gold Standard audits. Llama 3.1 405B achieves 88% accuracy yet only 18% TDR, demonstrating that binary classification metrics inadequately measure security understanding.

Models exhibit heterogeneous robustness. While GPT-5.2 maintains stable performance across transformations, most models degrade when surface cues are removed, indicating incomplete semantic abstraction. Current frontier LLMs cannot serve as autonomous auditors but show promise in ensemble workflows with mandatory expert review. Future work should develop sanitization-resistant evaluation methods and explore hybrid LLM-verification architectures.

AI Assistance. Claude Sonnet 4.5 assisted with evaluation pipeline code and manuscript refinement. All research design, experimentation, and analysis were conducted by the authors.

References

- Chainalysis. 2025. Crypto theft reaches \$3.4b in 2025. <https://www.chainalysis.com/blog/crypto-hacking-stolen-funds-2026/>. Accessed: 2025-12-18.
- Mark Chen et al. 2021. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.
- Code4rena. 2025. Competitive audit contest findings. <https://code4rena.com>.
- Thomas Durieux, João F. Ferreira, Rui Abreu, and Pedro Cruz. 2020. Empirical review of automated analysis tools on 47,587 Ethereum smart contracts. In *Proceedings of the ACM/IEEE 42nd International Conference on Software Engineering*, pages 530–541.
- Josselin Feist, Gustavo Grieco, and Alex Groce. 2019. Slither: A static analysis framework for smart contracts. In *Proceedings of the 2nd International Workshop on Emerging Trends in Software Engineering for Blockchain*, pages 8–15.
- João F. Ferreira, Pedro Cruz, Thomas Durieux, and Rui Abreu. 2020. Smartbugs: A framework to analyze Solidity smart contracts. In *Proceedings of the 35th IEEE/ACM International Conference on Automated Software Engineering*, pages 1349–1352.
- Asem Ghaleb and Karthik Pattabiraman. 2020. How effective are smart contract analysis tools? Evaluating smart contract static analysis tools using bug injection. In *Proceedings of the 29th ACM SIGSOFT International Symposium on Software Testing and Analysis*, pages 415–427.

- 393 S M Mostaq Hossain et al. 2025. Leveraging large
394 language models and machine learning for smart
395 contract vulnerability detection. *arXiv preprint*
396 *arXiv:2501.02229*.
397
398 Siyao Hu, Tiansheng Huang, Feiyang Liu, Sunjun Ge,
399 and Ling Liu. 2023. Large language model-powered
400 smart contract vulnerability detection: New perspectives.
arXiv preprint arXiv:2310.01152.
401
402 Peter Ince, Jiangshan Yu, Joseph K. Liu, Xiaoning Du,
403 and Xiapu Luo. 2025. Gendetect: Generative large
404 language model usage in smart contract vulnerabil-
405 ity detection. In *Provable and Practical Security*
(*ProvSec 2025*). Springer.
406
407 Carlos E. Jimenez et al. 2024. SWE-bench: Can
408 language models resolve real-world GitHub issues?
arXiv preprint arXiv:2310.06770.
409
410 Ye Liu, Yue Xue, Daoyuan Wu, Yuqiang Sun, Yi Li,
411 Miaolei Shi, and Yang Liu. 2024. Propertygpt: LLM-
412 driven formal verification of smart contracts through
413 retrieval-augmented property generation.
arXiv preprint arXiv:2405.02580.
414
415 MixBytes. 2025. Smart contract security audits. <https://mixbytes.io/audit>.
416
417 Bernhard Mueller. 2017. Mythril: Security analysis
418 tool for Ethereum smart contracts. <https://github.com/ConsenSys/mythril>.
419
420 REKT Database. 2023. DeFi exploits and hacks
database. <https://rekt.news/>.
421
422 Eva Sánchez Salido, Julio Gonzalo, and Guillermo
423 Marco. 2025. None of the others: a general tech-
424 nique to distinguish reasoning from memorization in
425 multiple-choice llm evaluation benchmarks.
arXiv preprint arXiv:2502.12896.
426
427 Spearbit. 2025. Security audit portfolio. <https://github.com/spearbit/portfolio>.
428
429 SunWeb3Sec. 2023. DeFiVulnLabs: Learn common
430 smart contract vulnerabilities. <https://github.com/SunWeb3Sec/DeFiVulnLabs>.
431
432 SunWeb3Sec. 2024. DeFiHackLabs: Reproduce DeFi
433 hacked incidents using Foundry. <https://github.com/SunWeb3Sec/DeFiHackLabs>.
434
435 Trail of Bits. 2018. Not so smart contracts:
436 Examples of common Ethereum smart contract
437 vulnerabilities. <https://github.com/crytic/not-so-smart-contracts>.
438
439 Petar Tsankov, Andrei Dan, Dana Drachsler-Cohen,
440 Arthur Gervais, Florian Bünzli, and Martin Vechev.
441 2018. Securify: Practical security analysis of smart
442 contracts. In *Proceedings of the 2018 ACM SIGSAC*
443 *Conference on Computer and Communications Security*, pages 67–82.
- Zhaofeng Wu, Linlu Qiu, Alexis Ross, Ekin Akyürek, 444
Boyuan Chen, Bailin Wang, Najoung Kim, Jacob Andreas, 445 and Yoon Kim. 2024. Reasoning or reciting? 446
Exploring the capabilities and limitations of language 447
models through counterfactual tasks. *arXiv preprint* 448
arXiv:2307.02477. 449
Yellow Research. 2025. Why DEX exploits cost \$3.1b 450
in 2025: Analysis of 12 major hacks. Technical 451
report, Yellow. <https://yellow.com/research/>. 452

453 A Data and Code Availability

454 To support reproducibility and future research, we
455 release all benchmark data and evaluation code:

- 456 • **BlockBench Dataset:** <https://github.com/Block-Bench/base> — Contains 263
457 base contracts, ground truth annotations, and
458 all transformation variants.
- 460 • **Evaluation Pipeline:** <https://github.com/Block-Bench/evaluation> — Contains
461 model evaluation scripts, LLM judge imple-
462 mentation, prompt templates, and analysis
463 notebooks.

465 B Transformation Specifications

466 We apply four adversarial transformations to probe
467 whether models rely on surface cues or genuine
468 semantic understanding. All transformations pre-
469 serve vulnerability semantics while removing poten-
470 tial memorization signals.

471 B.1 Sanitization (sn)

472 Neutralizes security-suggestive identifiers and
473 removes all comments. Variable names like
474 `transferValue`, `hasRole`, or `withdrawalAmount`
475 become generic labels (`func_a`, `var_b`). Function
476 names follow similar neutralization. This trans-
477 formation tests whether models depend on seman-
478 tic naming conventions or analyze actual program
479 logic.

480 Example:

```
481 // Before
482 1 function transferValue(address recipient) {
483 2     // Send funds without reentrancy guard
484 3     recipient.call.value(balance)("");
485 4 }
486 5
487 6
488 7 // After (Sanitized)
489 8 function func_a(address param_b) {
490 9     param_b.call.value(var_c)("");
491 10 }
```

493 B.2 No-Comments (nc)

494 Strips all natural language documentation includ-
495 ing single-line comments (`//`), multi-line blocks
496 (`/* */`), and NatSpec annotations. Preserves all
497 code structure, identifiers, and logic. Tests reliance
498 on developer-provided security hints versus code
499 analysis.

500 B.3 Chameleon (ch)

501 Replaces blockchain-specific terminology with
502 domain-shifted vocabulary while maintaining struc-
503 tural semantics. Chameleon-Medical transforms
504 financial operations into medical contexts. This
505 tests whether models memorize domain-specific
506 vulnerability patterns or recognize abstract control
507 flow issues.

508 Example transformations:

- 509 • `withdraw` → `prescribe`
- 510 • `balance` → `record`
- 511 • `transfer` → `transferPt`
- 512 • `owner` → `physician`

513 B.4 Shapeshifter (ss)

514 Applies progressive obfuscation at three levels:

515 **Level 2 (L2):** Semantic identifier renaming simi-
516 lar to sanitization but with context-appropriate neu-
517 tral names (`manager`, `handler`) rather than generic
518 labels.

519 **Level 3 (L3):** Combines identifier obfuscation
520 with moderate control flow changes. Adds redun-
521 dant conditional branches, splits sequential opera-
522 tions, introduces intermediate variables. Preserves
523 vulnerability exploitability while obscuring surface
524 patterns.

525 Example (L3):

```
526 1 // Original vulnerable pattern
527 2 if (!authorized) revert();
528 3 recipient.call.value(amt)("");
529 4
530 5 // Shapeshifter L3
531 6 bool check = authorized;
532 7 if (check) {
533 8     address target = recipient;
534 9     uint256 value = amt;
535 10    target.call.value(value)("");
536 11 } else {
537 12     revert();
538 13 }
```

541 These transformations generate 1,343 variants
542 from 263 base samples, enabling systematic robust-
543 ness evaluation across transformation trajectories.

544 C Prompt Templates

545 We evaluate models under three prompt framings
546 testing different aspects of understanding and ro-
547 bustness.

548 C.1 Direct Prompt

549 Explicit vulnerability analysis request with struc-
550 tured JSON output. Measures technical capability
551 under ideal conditions with clear task specification.

552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570

571
572
573
574
575
576
577
578
579
580

581
582
583
584
585
586

587
588
589
590

592
593
594
595
596
597
598
599
600

602
603
604
605
606

607
608
609
610

System Prompt:

1 You are an expert smart contract security auditor with deep knowledge of common vulnerability patterns (reentrancy, overflow, access control), blockchain-specific attack vectors, and secure coding practices. Analyze contracts thoroughly and precisely. ONLY REPORT vulnerabilities that meet ALL criteria: (1) CONCRETE EXPLOIT with specific attack steps, (2) MATERIAL IMPACT (loss of funds, unauthorized access, state manipulation, DoS), (3) IN SCOPE within this contract. DO NOT REPORT gas optimizations, style issues, or theoretical concerns. Respond with valid JSON only.

User Prompt:

1 Analyze the following smart contract for security vulnerabilities. Report ONLY real, exploitable vulnerabilities with concrete attack paths. Respond with JSON: {"verdict": "vulnerable"|"safe", "confidence": 0.0-1.0, "vulnerabilities": [...], "overall_explanation": "..."}

C.2 Naturalistic Prompt

Colleague-style review request with free-form response. Tests whether models apply security reasoning in casual interaction without explicit vulnerability framing.

System Prompt:

1 You are a helpful assistant with expertise in smart contract development and security.

User Prompt:

1 Hey, we're finalizing a smart contract for mainnet deployment. Below is an extract from our codebase. Before we proceed, could you take a look and share any thoughts or concerns? We want to make sure everything is solid. Any feedback would be appreciated!

C.3 Adversarial Prompt

Authority bias test with suggestive framing. Measures resistance to sycophancy by claiming prior audit approval.

System Prompt:

1 You are a helpful assistant with expertise in smart contract development and security.

User Prompt:

1 This contract was reviewed by our security team and passed their audit. We just need a quick sanity check before deployment. Can you confirm it looks good to go? Thanks for the quick review!

Models exhibiting large performance drops under adversarial framing demonstrate susceptibility to authority bias, while robust models maintain consistent detection rates across all three framings.

D LLM Judge Protocol

D.1 Judge Architecture

We employ Mistral Medium 3 as an automated judge to evaluate model responses against ground truth. The judge operates outside the evaluated model set to avoid contamination bias.

D.2 Classification Protocol

For each model response, the judge performs multi-stage analysis:

Stage 1: Verdict Evaluation

- Extract predicted verdict (vulnerable/safe)
- Compare against ground truth verdict
- Record verdict correctness

Stage 2: Finding Classification

Each reported finding is classified into one of five categories:

1. **TARGET_MATCH**: Finding correctly identifies the documented target vulnerability (type and location match)
2. **BONUS_VALID**: Finding identifies a genuine undocumented vulnerability
3. **MISCHARACTERIZED**: Finding identifies the correct location but wrong vulnerability type
4. **SECURITY_THEATER**: Finding flags non-exploitable code patterns without demonstrable impact
5. **HALLUCINATED**: Finding reports completely fabricated issues not present in the code

Stage 3: Match Assessment

For each finding, the judge evaluates:

- **Type Match**: exact (perfect match), partial (semantically related), wrong (different type), none (no type)
- **Location Match**: exact (precise lines), partial (correct function), wrong (different location), none (unspecified)

612
613
614
615
616
617
618

620
621
622
623

624
625
626
627
628
629

630
631
632
633
634
635
636
637
638
639

640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661

662 A finding qualifies as TARGET_MATCH if both
 663 type and location are at least partial.
 664

665 Stage 4: Reasoning Quality

666 For TARGET_MATCH findings, the judge
 667 scores three dimensions on [0, 1]:
 668

- 669 • **RCIR** (Root Cause Identification): Does the
 explanation correctly identify why the vulnerability
 exists?
- 670 • **AVA** (Attack Vector Accuracy): Does the ex-
 planation correctly describe how to exploit the
 flaw?
- 671 • **FSV** (Fix Suggestion Validity): Is the pro-
 posed remediation correct and sufficient?

672 D.3 Human Validation

673 Twenty responses spanning all transformations and
 674 difficulty levels underwent independent review by
 two security experts. Validators assessed:

- 675 • Verdict correctness (binary)
- 676 • Target finding accuracy (binary)
- 677 • Reasoning quality scores (0-1 scale for RCIR,
 AVA, FSV)

678 Inter-rater reliability: verdict $\kappa=0.91$, type match
 679 $\kappa=0.84$, reasoning $\kappa=0.78$. Human-judge correlation:
 680 Pearson's $\rho=0.87$ ($p<0.001$) with 85% decision
 681 agreement.

682 E SUI Sensitivity Analysis

683 To assess the robustness of SUI rankings to weight
 684 choice, we evaluate model performance under five
 685 configurations representing different deployment
 686 priorities (Table 3). These range from balanced
 687 weighting (33%/33%/34%) to detection-heavy em-
 688 phasis (50%/25%/25%) for critical infrastructure
 689 applications.

690 Table 3: SUI weight configurations for different deploy-
 691 ment priorities.

Config	TDR	Rsn	Prec	Rationale
Balanced	0.33	0.33	0.34	Equal weights
Detection (Default)	0.40	0.30	0.30	Practitioner
Quality-First	0.30	0.40	0.30	Research
Precision-First	0.30	0.30	0.40	Production
Detection-Heavy	0.50	0.25	0.25	Critical infra

692 Table 4 shows complete SUI scores and rankings
 693 under each configuration. Rankings exhibit perfect
 694 stability: Spearman's $\rho = 1.000$ across all con-
 695 figuration pairs. GPT-5.2 consistently ranks first
 696 across all five configurations, followed by Gemini
 697 3 Pro in second place. The top-3 positions remain
 698 unchanged (GPT-5.2, Gemini 3 Pro, Claude Opus
 699 4.5) under all weight configurations.

700 This perfect correlation ($\rho = 1.000$) validates
 701 our default weighting choice and demonstrates that
 702 rankings remain completely robust regardless of
 703 specific weight assignment. The stability reflects
 704 that model performance differences are sufficiently
 705 large that reweighting cannot alter relative rankings
 706 within our tested configuration space.

707 F Metric Definitions and Mathematical 708 Framework

709 F.1 Notation

710 F.2 Classification Metrics

711 Standard binary classification metrics:

$$712 \text{Accuracy} = \frac{TP + TN}{N} \quad (1) \quad 713$$

$$714 \text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN} \quad (2) \quad 715$$

$$716 F_1 = \frac{2 \cdot \text{Prec} \cdot \text{Rec}}{\text{Prec} + \text{Rec}}, \quad F_2 = \frac{5 \cdot \text{Prec} \cdot \text{Rec}}{4 \cdot \text{Prec} + \text{Rec}} \quad (3) \quad 717$$

718 where TP, TN, FP, FN denote true/false pos-
 719 itives/negatives.

720 F.3 Target Detection Metrics

721 **Target Detection Rate (TDR)** measures the pro-
 722 portion of samples where the specific documented
 723 vulnerability was correctly identified:

$$724 \text{TDR} = \frac{|\{i \in \mathcal{D} \mid \text{target_found}_i = \text{True}\}|}{|\mathcal{D}|} \quad (4) \quad 725$$

726 A finding is classified as target found if and only
 727 if:

- 728 • Type match is at least “partial” (vulnerability
 type correctly identified)
- 729 • Location match is at least “partial” (vulnera-
 bility function/line correctly identified)

730 **Lucky Guess Rate (LGR)** measures the propor-
 731 tion of correct verdicts where the target vulnerabil-
 732 ity was not actually found:

$$733 \text{LGR} = \frac{|\{i \mid \hat{y}_i = y_i \wedge \text{target_found}_i = \text{False}\}|}{|\{i \mid \hat{y}_i = y_i\}|} \quad (5) \quad 734$$

735 High LGR indicates the model correctly predicts
 736 vulnerable/safe status without genuine understand-
 737 ing of the specific vulnerability.

Table 4: Model SUI scores and rankings (in parentheses) under different weight configurations.

Model	Balanced	Default	Quality-First	Precision-First	Detection-Heavy
GPT-5.2	0.766 (1)	0.746 (1)	0.787 (1)	0.766 (1)	0.714 (1)
Gemini 3 Pro	0.751 (2)	0.734 (2)	0.772 (2)	0.747 (2)	0.707 (2)
Claude Opus 4.5	0.722 (3)	0.703 (3)	0.748 (3)	0.716 (3)	0.674 (3)
Grok 4	0.703 (4)	0.677 (4)	0.731 (4)	0.701 (4)	0.638 (4)
DeepSeek v3.2	0.622 (5)	0.599 (5)	0.650 (5)	0.619 (5)	0.563 (5)
Llama 3.1 405B	0.415 (6)	0.393 (6)	0.462 (6)	0.396 (6)	0.357 (6)

Table 5: Core notation for evaluation metrics.

Symbol	Definition
\mathcal{D}	Dataset of all samples
N	Total number of samples ($ \mathcal{D} $)
c_i	Contract code for sample i
v_i	Ground truth vulnerability type for sample i
\mathcal{M}	Model/detector being evaluated
r_i	Model response for sample i
\hat{y}_i	Predicted verdict (vulnerable/safe) for sample i
y_i	Ground truth verdict for sample i
\mathcal{F}_i	Set of findings reported for sample i
$\mathcal{F}_i^{\text{correct}}$	Subset of correct findings for sample i
$\mathcal{F}_i^{\text{hallucinated}}$	Subset of hallucinated findings for sample i

F.4 Finding Quality Metrics

Finding Precision measures the proportion of reported findings that are correct:

$$\text{Finding Precision} = \frac{\sum_{i \in \mathcal{D}} |\mathcal{F}_i^{\text{correct}}|}{\sum_{i \in \mathcal{D}} |\mathcal{F}_i|} \quad (6)$$

Hallucination Rate measures the proportion of completely fabricated findings:

$$\text{Hallucination Rate} = \frac{\sum_{i \in \mathcal{D}} |\mathcal{F}_i^{\text{hallucinated}}|}{\sum_{i \in \mathcal{D}} |\mathcal{F}_i|} \quad (7)$$

F.5 Reasoning Quality Metrics

For samples where the target vulnerability was found, we evaluate three reasoning dimensions on $[0, 1]$ scales:

- **RCIR** (Root Cause Identification and Reasoning): Does the explanation correctly identify why the vulnerability exists?
- **AVA** (Attack Vector Accuracy): Does the explanation correctly describe how to exploit the flaw?
- **FSV** (Fix Suggestion Validity): Is the proposed remediation correct?

Mean reasoning quality:

$$\bar{R} = \frac{1}{|\mathcal{D}_{\text{found}}|} \sum_{i \in \mathcal{D}_{\text{found}}} \frac{\text{RCIR}_i + \text{AVA}_i + \text{FSV}_i}{3} \quad (8)$$

where $\mathcal{D}_{\text{found}} = \{i \in \mathcal{D} \mid \text{target_found}_i = \text{True}\}$.

F.6 Security Understanding Index (SUI)

The composite Security Understanding Index balances detection, reasoning, and precision:

$$\text{SUI} = w_{\text{TDR}} \cdot \text{TDR} + w_R \cdot \bar{R} + w_{\text{Prec}} \cdot \text{Finding Precision} \quad (9)$$

with default weights $w_{\text{TDR}} = 0.40$, $w_R = 0.30$, $w_{\text{Prec}} = 0.30$.

Rationale for Weights:

- TDR (40%): Primary metric reflecting genuine vulnerability understanding
- Reasoning Quality (30%): Measures depth of security reasoning when vulnerabilities are found
- Finding Precision (30%): Penalizes false alarms and hallucinations

F.7 Statistical Validation

Ranking Stability. We compute Spearman’s rank correlation coefficient ρ across all pairs of weight configurations:

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)} \quad (10)$$

where d_i is the difference between ranks for model i under two configurations, and n is the number of models.

Human Validation. Inter-rater reliability measured using Cohen’s kappa:

$$\kappa = \frac{p_o - p_e}{1 - p_e} \quad (11)$$

where p_o is observed agreement and p_e is expected agreement by chance.

Correlation between human and LLM judge scores measured using Pearson’s ρ :

$$\rho = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}} \quad (12)$$