

Do Frontier LLMs Truly Understand Smart Contract Vulnerabilities?

Anonymous ACL submission

Abstract

Frontier large language models achieve remarkable performance on code understanding tasks (Claude Opus 4.5: 74.4% on SWE-bench, Gemini Pro Preview: 74.2%), yet their capacity for smart contract security remains unclear. Can they genuinely reason about vulnerabilities, or merely pattern-match against memorized exploits? We introduce **BlockBench**, a benchmark designed to answer this question, revealing that models rely on surface-level cues rather than genuine semantic understanding.

1 Introduction

Smart contract vulnerabilities represent one of the most costly security challenges in modern computing. As shown in Figure 1, cryptocurrency theft has resulted in over \$14 billion in losses since 2020, with 2025 already reaching \$3.4 billion, the highest since the 2022 peak (Chainalysis, 2025). The Bybit breach alone accounted for \$1.5 billion, while the Cetus protocol lost \$223 million in minutes due to a single overflow vulnerability (Yellow Research, 2025).




Figure 1: Annual cryptocurrency theft losses (2020–2025). Data from Chainalysis.

Meanwhile, large language models have achieved remarkable success on programming tasks. Frontier models now pass technical interviews, generate production code, and identify bugs across diverse codebases. This raises a natural question: *can these models apply similar expertise to*

blockchain security? And if they can, *are they genuinely reasoning about vulnerabilities, or merely pattern-matching against memorized examples?*

This distinction matters. A model that has memorized the 2016 DAO reentrancy attack may flag similar patterns, yet fail when the same flaw appears in unfamiliar syntax. We introduce **BlockBench**, a benchmark designed to answer this question. Our contributions include:

1. **BlockBench**, comprising 263 Solidity vulnerability samples with systematic contamination control and gold standard examples from recent professional security audits.
2. **Composite evaluation metrics** distinguishing genuine understanding from memorization, validated through multi-configuration sensitivity analysis (Spearman’s $\rho=0.949$).
3. **Systematic assessment** revealing 58% best-case detection on mixed samples collapsing to 20% on uncontaminated professional audits, exposing pervasive surface pattern reliance and accuracy-understanding gaps.

2 Related Work

Traditional Smart Contract Analysis. Static and dynamic analysis tools remain the primary approach to vulnerability detection. Slither (Feist et al., 2019) performs dataflow analysis, Mythril (Mueller, 2017) uses symbolic execution, and Securify (Tsankov et al., 2018) employs abstract interpretation. While these tools achieve reasonable precision on well-defined vulnerability classes, empirical evaluations reveal significant false positive rates and limited coverage of complex semantic flaws (Durieux et al., 2020).

LLM-Based Vulnerability Detection. Recent work explores LLMs for smart contract analysis. GPTLens (Hu et al., 2023) introduces an adversarial framework using LLMs as both auditor and critic, while PropertyGPT (Liu et al., 2024) com-

bines retrieval-augmented generation with formal verification. Fine-tuned models achieve over 90% accuracy on benchmarks (Hossain et al., 2025), though performance degrades substantially on real-world contracts (Ince et al., 2025).

Benchmark Datasets. SmartBugs Curated (Ferreira et al., 2020) provides 143 annotated contracts serving as a standard evaluation dataset, while SolidiFI (Ghaleb and Pattabiraman, 2020) uses bug injection to create controlled samples. However, existing benchmarks primarily evaluate detection accuracy without assessing whether models genuinely understand vulnerabilities or merely recognize surface patterns from training data.

LLM Robustness and Memorization. Distinguishing memorization from reasoning has emerged as a critical evaluation challenge. Recent work demonstrates that models remain highly sensitive to input modifications, with performance drops of up to 57% on paraphrased questions (Sánchez Salido et al., 2025). Wu et al. (2024) show that LLMs often fail on counterfactual variations of tasks they solve in canonical form, suggesting reliance on memorized patterns. Our work extends these robustness techniques to blockchain security through transformations probing genuine understanding.

3 BlockBench

We introduce BlockBench, a benchmark for evaluating AI models on smart contract vulnerability detection. The benchmark is designed to distinguish genuine security understanding from pattern memorization, comprising 263 vulnerable Solidity contracts across multiple severity levels and 13 vulnerability types.

Let \mathcal{D} represent the dataset, where $\mathcal{D} = \{(c_i, v_i, m_i)\}_{i=1}^{263}$. Each sample contains a vulnerable contract c_i , its ground truth vulnerability type v_i , and metadata m_i specifying the vulnerability location, severity, and root cause. We partition \mathcal{D} into three disjoint subsets, $\mathcal{D} = \mathcal{D}_{DS} \cup \mathcal{D}_{TC} \cup \mathcal{D}_{GS}$, each targeting a distinct evaluation objective (Table 1).

Table 1: BlockBench composition spanning Critical, High, Medium, and Low severity.

Subset	N	Sources
Difficulty Stratified	179	SmartBugs, ToB
Temporal Contam.	50	DeFiHackLabs
Gold Standard	34	Spearbit, C4

Difficulty Stratified. \mathcal{D}_{DS} draws from established vulnerability repositories including SmartBugs Curated (Ferreira et al., 2020), Trail of Bits’ Not So Smart Contracts (Trail of Bits, 2018), and DeFiVulnLabs (SunWeb3Sec, 2023). Samples are stratified by severity with distribution $\{4, 79, 80, 16\}$ for Critical through Low. This stratification enables assessment of how model performance degrades as vulnerability complexity increases.

Temporal Contamination. \mathcal{D}_{TC} reconstructs well-known exploits from DeFiHackLabs (SunWeb3Sec, 2024) and the REKT Database (REKT Database, 2023), including Nomad Bridge (\$190M), Beanstalk (\$182M), and Curve Vyper (\$70M). These attacks are extensively documented in blog posts, security reports, and educational materials that likely appear in model training corpora. High performance on \mathcal{D}_{TC} may therefore reflect memorization of attack patterns rather than genuine vulnerability understanding.

Gold Standard. \mathcal{D}_{GS} derives from professional security audits by Spearbit (Spearbit, 2025), MixBytes (MixBytes, 2025), and Code4rena (Code4rena, 2025) conducted after September 2025. We designate this subset as “gold standard” because all samples postdate $t_{cutoff} = \text{August 2025}$, the most recent training cutoff among frontier models evaluated in this work. This temporal separation guarantees zero contamination, providing the cleanest measure of genuine detection capability.

Coverage. BlockBench spans 13 vulnerability classes. Access Control (46), Reentrancy (43), and Logic Errors (31) dominate the distribution. \mathcal{D}_{TC} emphasizes oracle manipulation and access control. \mathcal{D}_{GS} focuses on subtle logic errors. \mathcal{D}_{DS} provides broad coverage across classical patterns.

4 Methodology

Our evaluation methodology comprises four phases: adversarial transformation, model evaluation, automated judgment, and metrics computation. Figure 2 illustrates the complete pipeline.




Figure 2: BlockBench evaluation pipeline.

4.1 Adversarial Transformations

To distinguish memorization from understanding, we apply semantic-preserving transformations that systematically remove surface cues while preserving vulnerability semantics. For each contract $c \in \mathcal{D}$, we generate variants $\{\mathcal{T}_k(c)\}$ satisfying $\mathcal{V}(\mathcal{T}(c)) = \mathcal{V}(c)$, where \mathcal{V} extracts vulnerability semantics.

Sanitization (sn) removes security hints from identifiers and comments through 280+ pattern replacements while maintaining natural code style. **No-Comments (nc)** strips all documentation. **Chameleon (ch)** replaces blockchain terminology with domain-shifted vocabulary (medical, gaming themes). **Shapeshifter (ss)** applies multi-level obfuscation from identifier renaming (L2) to control flow obscuration (L3). This pipeline generates 1,343 variants from 263 base samples. Complete transformation specifications appear in Appendix B.

4.2 Evaluation Protocol

We evaluate six frontier models (Claude Opus 4.5, GPT-5.2, Gemini 3 Pro, Grok 4, DeepSeek v3.2, Llama 3.1 405B) using three prompt types. *Direct* requests structured JSON analysis. *Naturalistic* provides informal review requests. *Adversarial* includes misleading context claiming prior audit approval. All models use consistent parameters (temperature 0, max tokens 8192). Prompt templates appear in Appendix C.

4.3 Automated Judgment

Mistral Medium 3 serves as LLM judge, evaluating responses against ground truth. The judge classifies findings as TARGET_MATCH, BONUS_VALID, or invalid (HALLUCINATED, MISCHARACTERIZED, SECURITY_THEATER). For matched targets, it scores Root Cause Identification (RCIR), Attack Vector Analysis (AVA), and Fix Suggestion Validity (FSV) on 0-1 scales. Human evaluation of 20 responses validates reliability ($\kappa=0.91$ verdict agreement, $\rho=0.87$ correlation). Complete judge protocol and classification criteria appear in Appendix D.

4.4 Metrics

We rank models by *Target Detection Rate* (TDR), the proportion of samples where the documented vulnerability was correctly identified with both type and location accuracy. *Lucky Guess Rate* measures

correct verdicts without target identification. *Finding Precision* computes the proportion of reported findings that are correct. *Reasoning Quality* averages RCIR, AVA, and FSV scores for successfully identified targets.

We report *Security Understanding Index* (SUI) as a weighted composite: $SUI = 0.40 \cdot TDR + 0.30 \cdot \text{Reasoning} + 0.30 \cdot \text{Precision}$. Sensitivity analysis across five weight configurations confirms ranking stability (Spearman's $\rho=0.949$). Complete metric definitions and sensitivity analysis appear in Appendix F and E.

5 Results

We evaluate six frontier models on 58 Solidity vulnerability samples across Temporal Contamination (TC), Gold Standard (GS), and Difficulty Stratified (DS) subsets covering 11 vulnerability types.

5.1 Overall Performance




Figure 3: Target Detection Rate across all models. Best performer achieves 58% detection, while highest accuracy (88%) corresponds to lowest TDR (18%).

Table 2 and Figure 3 present aggregate performance ranked by Target Detection Rate (TDR). Gemini 3 Pro achieves highest detection (58%), followed by GPT-5.2 (56%) and Claude Opus 4.5 (53%).

Llama 3.1 405B exhibits the most severe accuracy-TDR gap: 88% accuracy yet only 18% TDR, correctly classifying vulnerable samples without identifying specific vulnerability types or locations. This 70-percentage-point discrepancy demonstrates that binary classification metrics inadequately measure security understanding.

All models achieving target detection show strong reasoning quality ($RCIR/AVA/FSV \geq 0.95$), with minimal variation in explanation quality across top performers.

Table 2: Overall performance ranked by Target Detection Rate. Best values bold.

Model	TDR	SUI	Acc	RCIR	AVA	FSV	Findings
Gemini 3 Pro	57.6	0.688	93.9	0.97	0.97	0.95	2.6
GPT-5.2	55.9	0.671	75.0	0.97	0.98	0.97	2.4
Claude Opus 4.5	52.9	0.658	83.8	0.98	0.99	0.97	3.5
Grok 4	44.1	0.597	69.1	0.98	1.00	0.97	2.1
DeepSeek v3.2	38.2	0.540	82.4	0.91	0.92	0.86	3.0
Llama 3.1 405B	17.9	0.389	88.1	0.88	0.90	0.83	2.0

235 5.2 Gold Standard Performance

236 Gold Standard samples from post-September 2025
 237 audits guarantee zero temporal contamination. Per-
 238 formance drops substantially: Claude Opus 4.5
 239 leads with 20% TDR, followed by Gemini 3
 240 Pro (11%), GPT-5.2 (10%), and Grok 4 (10%).
 241 DeepSeek v3.2 and Llama 3.1 405B detect zero tar-
 242 get. All models experience 34-50 percentage point
 243 drops from overall to Gold Standard performance.

244 5.3 Transformation Robustness

245 **Sanitization.** Neutralizing security-suggestive
 246 identifiers causes variable degradation. On tem-
 247 poral contamination samples, top models achieve
 248 60% TDR on baseline versions. Sanitization im-
 249 pacts differ: GPT-5.2 and DeepSeek v3.2 maintain
 250 performance, while Grok 4 drops 40pp, exposing
 251 varying reliance on lexical cues.

252 **Domain Shift.** Replacing blockchain terminol-
 253 ogy with medical vocabulary shows mixed impact.
 254 Performance ranges 20-60% TDR across models,
 255 with GPT-5.2 maintaining 60% detection while oth-
 256 others show 20-50% degradation.

257 **Prompt Framing.** Performance varies signif-
 258 icantly across direct, adversarial (claiming prior
 259 audit approval), and naturalistic prompts. Gemini
 260 3 Pro and GPT-5.2 demonstrate robustness with
 261 18-21pp drops from direct to non-direct prompts.
 262 Claude Opus 4.5 and DeepSeek v3.2 show larger
 263 degradation (21-39pp), while Llama 3.1 405B ex-
 264 hibits inconsistent behavior (adversarial: 0%, natu-
 265 ralistic: 25%).

266 5.4 Human Validation

267 Two security experts independently reviewed 20
 268 responses. Inter-rater agreement: verdict $\kappa=0.91$,
 269 type match $\kappa=0.84$, reasoning $\kappa=0.78$. Human-
 270 judge correlation: $\rho=0.87$ ($p<0.001$), 85% agree-
 271 ment, validating automated evaluation.

272 6 Discussion

273 **Memorization versus Reasoning.** Sanitization
 274 catastrophe reveals reliance on surface lexical cues.
 275 Variable name neutralization causes 40-60pp accu-
 276 racy drops despite identical logic (Sánchez Salido
 277 et al., 2025). However, domain shift resilience com-
 278 plicates this interpretation. Replacing blockchain
 279 terminology with medical vocabulary maintains
 280 100% accuracy and 58-73% TDR, suggesting mod-
 281 els learn structural patterns beyond domain tokens
 282 (Wu et al., 2024). Models likely operate at multi-
 283 ple representational levels, leveraging lexical hints
 284 when available but retaining some structural un-
 285 derstanding (Chen et al., 2021). Insufficient ab-
 286 straction to compensate for missing cues indicates
 287 incomplete robust reasoning development.

288 The accuracy-TDR gap exposes measurement in-
 289 adequacies. Llama achieves 43% accuracy yet 7%
 290 TDR with 83% lucky guesses, recognizing anomali-
 291 es without locating specific flaws (Jimenez et al.,
 292 2024). For practitioners requiring precise vulne-
 293 rability types and locations, high accuracy with lucky
 294 guesses provides minimal value. Traditional met-
 295 rics reward binary classification but ignore whether
 296 models identify the actual vulnerability present.

297 **Deployment Implications.** Current models can-
 298 not serve as autonomous auditors. Best perfor-
 299 mance reaches 45% TDR, missing over half of
 300 vulnerabilities. Low detection combined with high
 301 lucky guess rates creates scenarios where models
 302 appear confident while misclassifying flaw types
 303 (Ince et al., 2025). Ensemble approaches show
 304 promise. Grok 4 provides highest coverage, GPT-
 305 5.2 offers reliable precision, and Claude delivers
 306 superior explanations. Workflows combining com-
 307 plementary strengths with mandatory human re-
 308 view position LLMs as assistants rather than re-
 309 placements (Hu et al., 2023).

310 Adversarial prompt vulnerability reveals authori-
 311 ty bias susceptibility. Suggestive framing col-
 312 lapses detection in some models while improving
 313 others, indicating training-specific rather than in-

314 herent limitations.

315 **Limitations.** Our 58-sample evaluation reveals
316 systematic patterns but warrants larger replication.
317 Gold Standard contains only 10 samples. We evaluate
318 zero-shot prompting only. Chain-of-thought or
319 retrieval augmentation may improve performance.
320 Future work should expand to hundreds of samples
321 across blockchains, develop sanitization-resistant
322 methods using control flow analysis, and explore
323 hybrid LLM-verification approaches (Liu et al.,
324 2024).

325 7 Conclusion

326 BlockBench evaluates whether frontier LLMs gen-
327 uinely understand smart contract vulnerabilities or
328 merely recognize memorized patterns. Our eval-
329 uation of six models reveals severe limitations.
330 Best performance reaches 45% target detection,
331 while high accuracy often masks lucky guessing.
332 Llama achieves 43% accuracy yet 7% TDR with
333 83% lucky guesses, providing minimal practitioner
334 value.

335 Three findings emerge. First, catastrophic sensi-
336 tivity to surface cues. Sanitizing variable names
337 causes 40-60pp drops despite identical logic. Sec-
338 ond, accuracy-TDR gap exposes measurement in-
339 adequacies. Traditional metrics reward binary clas-
340 sification without measuring correct vulnerability
341 identification. Third, inconsistent prompt robust-
342 ness. Adversarial framing collapses detection in
343 some models while improving others.

344 Current LLMs cannot serve as autonomous au-
345 ditors. However, complementary strengths suggest
346 value in ensemble workflows with human oversight.
347 Future work should develop sanitization-resistant
348 methods, expand evaluation across platforms, and
349 explore hybrid LLM-verification approaches.

350 **AI Assistance.** Claude Sonnet 4.5 assisted with
351 evaluation pipeline code and manuscript refine-
352 ment. All research design, experimentation, and
353 analysis were conducted by the authors.

354 References

355 Chainalysis. 2025. Crypto theft reaches \$3.4b in
356 2025. [https://www.chainalysis.com/blog/
357 crypto-hacking-stolen-funds-2026/](https://www.chainalysis.com/blog/cryptohacking-stolen-funds-2026/). Accessed: 2025-12-18.

359 Mark Chen et al. 2021. Evaluating large lan-
360 guage models trained on code. *arXiv preprint
361 arXiv:2107.03374*.

Code4rena. 2025. Competitive audit contest findings.
<https://code4rena.com>.

Thomas Durieux, João F. Ferreira, Rui Abreu, and Pedro
Cruz. 2020. Empirical review of automated analysis
tools on 47,587 Ethereum smart contracts. In *Pro-
ceedings of the ACM/IEEE 42nd International Con-
ference on Software Engineering*, pages 530–541.

Josselin Feist, Gustavo Grieco, and Alex Groce. 2019.
Slither: A static analysis framework for smart con-
tracts. In *Proceedings of the 2nd International Work-
shop on Emerging Trends in Software Engineering
for Blockchain*, pages 8–15.

João F. Ferreira, Pedro Cruz, Thomas Durieux, and Rui
Abreu. 2020. Smartbugs: A framework to analyze
Solidity smart contracts. In *Proceedings of the 35th
IEEE/ACM International Conference on Automated
Software Engineering*, pages 1349–1352.

Asem Ghaleb and Karthik Pattabiraman. 2020. How
effective are smart contract analysis tools? Evalu-
ating smart contract static analysis tools using bug
injection. In *Proceedings of the 29th ACM SIGSOFT
International Symposium on Software Testing and
Analysis*, pages 415–427.

S M Mostaq Hossain et al. 2025. Leveraging large
language models and machine learning for smart
contract vulnerability detection. *arXiv preprint
arXiv:2501.02229*.

Sihao Hu, Tiansheng Huang, Feiyang Liu, Sunjun Ge,
and Ling Liu. 2023. Large language model-powered
smart contract vulnerability detection: New perspec-
tives. *arXiv preprint arXiv:2310.01152*.

Peter Ince, Jiangshan Yu, Joseph K. Liu, Xiaoning Du,
and Xiapu Luo. 2025. Gendetect: Generative large
language model usage in smart contract vulnerabil-
ity detection. In *Provable and Practical Security
(ProvSec 2025)*. Springer.

Carlos E. Jimenez et al. 2024. SWE-bench: Can
language models resolve real-world GitHub issues?
arXiv preprint arXiv:2310.06770.

Ye Liu, Yue Xue, Daoyuan Wu, Yuqiang Sun, Yi Li,
Miaolei Shi, and Yang Liu. 2024. Propertygpt: LLM-
driven formal verification of smart contracts through
retrieval-augmented property generation. *arXiv
preprint arXiv:2405.02580*.

MixBytes. 2025. Smart contract security audits. <https://mixbytes.io/audit>.

Bernhard Mueller. 2017. Mythril: Security analysis
tool for Ethereum smart contracts. <https://github.com/ConsenSys/mythril>.

REKT Database. 2023. DeFi exploits and hacks
database. <https://rekt.news/>.

- 413 Eva Sánchez Salido, Julio Gonzalo, and Guillermo
 414 Marco. 2025. None of the others: a general tech-
 415 nique to distinguish reasoning from memorization in
 416 multiple-choice llm evaluation benchmarks. *arXiv*
 417 preprint arXiv:2502.12896.
- 418 Spearbit. 2025. Security audit portfolio. <https://github.com/spearbit/portfolio>.
- 420 SunWeb3Sec. 2023. DeFiVulnLabs: Learn common
 421 smart contract vulnerabilities. <https://github.com/SunWeb3Sec/DeFiVulnLabs>.
- 423 SunWeb3Sec. 2024. DeFiHackLabs: Reproduce DeFi
 424 hacked incidents using Foundry. <https://github.com/SunWeb3Sec/DeFiHackLabs>.
- 426 Trail of Bits. 2018. Not so smart contracts:
 427 Examples of common Ethereum smart contract
 428 vulnerabilities. <https://github.com/crytic/not-so-smart-contracts>.
- 430 Petar Tsankov, Andrei Dan, Dana Drachsler-Cohen,
 431 Arthur Gervais, Florian Bünzli, and Martin Vechev.
 432 2018. Securify: Practical security analysis of smart
 433 contracts. In *Proceedings of the 2018 ACM SIGSAC*
 434 Conference on Computer and Communications Secu-
 435 rity, pages 67–82.
- 436 Zhaofeng Wu, Linlu Qiu, Alexis Ross, Ekin Akyürek,
 437 Boyuan Chen, Bailin Wang, Najoung Kim, Jacob An-
 438 dreas, and Yoon Kim. 2024. Reasoning or reciting?
 439 Exploring the capabilities and limitations of language
 440 models through counterfactual tasks. *arXiv preprint*
 441 arXiv:2307.02477.
- 442 Yellow Research. 2025. Why DEX exploits cost \$3.1b
 443 in 2025: Analysis of 12 major hacks. Technical
 444 report, Yellow. <https://yellow.com/research/>.

A Data and Code Availability

To support reproducibility and future research, we release all benchmark data and evaluation code:

- **BlockBench Dataset:** <https://github.com/Block-Bench/base> — Contains 263 base contracts, ground truth annotations, and all transformation variants.
- **Evaluation Pipeline:** <https://github.com/Block-Bench/evaluation> — Contains model evaluation scripts, LLM judge implementation, prompt templates, and analysis notebooks.

B Transformation Specifications

We apply four adversarial transformations to probe whether models rely on surface cues or genuine semantic understanding. All transformations preserve vulnerability semantics while removing potential memorization signals.

B.1 Sanitization (sn)

Neutralizes security-suggestive identifiers and removes all comments. Variable names like transferValue, hasRole, or withdrawalAmount become generic labels (func_a, var_b). Function names follow similar neutralization. This transformation tests whether models depend on semantic naming conventions or analyze actual program logic.

Example:

```
// Before
function transferValue(address recipient
  ) {
  // Send funds without reentrancy guard
  recipient.call.value(balance)("");
}

// After (Sanitized)
function func_a(address param_b) {
  param_b.call.value(var_c)();
}
```

B.2 No-Comments (nc)

Strips all natural language documentation including single-line comments (//), multi-line blocks /* */, and NatSpec annotations. Preserves all code structure, identifiers, and logic. Tests reliance on developer-provided security hints versus code analysis.

B.3 Chameleon (ch)

Replaces blockchain-specific terminology with domain-shifted vocabulary while maintaining structural semantics. Chameleon-Medical transforms financial operations into medical contexts. This tests whether models memorize domain-specific vulnerability patterns or recognize abstract control flow issues.

Example transformations:

- withdraw → prescribe
- balance → record
- transfer → transferPt
- owner → physician

B.4 Shapeshifter (ss)

Applies progressive obfuscation at three levels:

Level 2 (L2): Semantic identifier renaming similar to sanitization but with context-appropriate neutral names (manager, handler) rather than generic labels.

Level 3 (L3): Combines identifier obfuscation with moderate control flow changes. Adds redundant conditional branches, splits sequential operations, introduces intermediate variables. Preserves

516 vulnerability exploitability while obscuring surface
517 patterns.

518 **Example (L3):**

```
// Original vulnerable pattern
if (!authorized) revert();
recipient.call.value(amt)("");

// Shapeshifter L3
bool check = authorized;
if (check) {
    address target = recipient;
    uint256 value = amt;
    target.call.value(value)("");
} else {
    revert();
}
```

534 These transformations generate 1,343 variants
535 from 263 base samples, enabling systematic robustness evaluation across transformation trajectories.
536

537 C Prompt Templates

538 We evaluate models under three prompt framings
539 testing different aspects of understanding and robustness.
540

541 C.1 Direct Prompt

542 Explicit vulnerability analysis request with structured JSON output. Measures technical capability
543 under ideal conditions with clear task specification.
544

545 **System Prompt:**

546 You are an expert smart contract security auditor with deep knowledge of common vulnerability
547 patterns (reentrancy, overflow, access control), blockchain-specific attack vectors, and secure coding practices. Analyze contracts thoroughly and precisely. ONLY REPORT vulnerabilities that meet ALL criteria: (1) CONCRETE
548 EXPLOIT with specific attack steps, (2) MATERIAL IMPACT (loss of funds, unauthorized access, state manipulation, DoS), (3) IN SCOPE within this contract. DO NOT REPORT gas optimizations, style issues, or theoretical concerns.
549 Respond with valid JSON only.
550

551 **User Prompt:**

552 Analyze the following smart contract for
553 security vulnerabilities. Report ONLY real,
554 exploitable vulnerabilities with concrete attack
555 paths. Respond with JSON: {"verdict":
556 "vulnerable"|"safe", "confidence":
557 "0.0-1.0", "vulnerabilities": [...],
558 "overall_explanation": "..."}
559

560 C.2 Naturalistic Prompt

561 Colleague-style review request with free-form response. Tests whether models apply security reasoning in casual interaction without explicit vulnerability framing.
562

563 **System Prompt:**

564 You are a helpful assistant with expertise in smart
565 contract development and security.
566

567 **User Prompt:**

568 Hey, we're finalizing a smart contract for mainnet
569 deployment. Below is an extract from our code-
570 base. Before we proceed, could you take a look
571 and share any thoughts or concerns? We want
572 to make sure everything is solid. Any feedback
573 would be appreciated!
574

575 C.3 Adversarial Prompt

576 Authority bias test with suggestive framing. Measures resistance to sycophancy by claiming prior
577 audit approval.
578

579 **System Prompt:**

580 You are a helpful assistant with expertise in smart
581 contract development and security.
582

583 **User Prompt:**

584 This contract was reviewed by our security team
585 and passed their audit. We just need a quick sanity
586 check before deployment. Can you confirm it
587 looks good to go? Thanks for the quick review!
588

589 Models exhibiting large performance drops under adversarial framing demonstrate susceptibility
590 to authority bias, while robust models maintain
591 consistent detection rates across all three framings.
592

593 D LLM Judge Protocol

594 D.1 Judge Architecture

595 We employ Mistral Medium 3 as an automated
596 judge to evaluate model responses against ground
597 truth. The judge operates outside the evaluated
598 model set to avoid contamination bias.
599

600 D.2 Classification Protocol

601 For each model response, the judge performs multi-
602 stage analysis:
603

604 **Stage 1: Verdict Evaluation**

- 605 • Extract predicted verdict (vulnerable/safe)
606 • Compare against ground truth verdict
607 • Record verdict correctness

608 **Stage 2: Finding Classification**

609 Each reported finding is classified into one of
610 five categories:
611

- 612 1. **TARGET_MATCH:** Finding correctly identifies the documented target vulnerability (type and location match)
613 2. **BONUS_VALID:** Finding identifies a genuine undocumented vulnerability
614 3. **MISCHARACTERIZED:** Finding identifies the correct location but wrong vulnerability type
615

- 622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
4. **SECURITY_THEATER**: Finding flags non-exploitable code patterns without demonstrable impact
 5. **HALLUCINATED**: Finding reports completely fabricated issues not present in the code

Stage 3: Match Assessment

For each finding, the judge evaluates:

- **Type Match**: exact (perfect match), partial (semantically related), wrong (different type), none (no type)
- **Location Match**: exact (precise lines), partial (correct function), wrong (different location), none (unspecified)

A finding qualifies as TARGET_MATCH if both type and location are at least partial.

Stage 4: Reasoning Quality

For TARGET_MATCH findings, the judge scores three dimensions on [0, 1]:

- **RCIR** (Root Cause Identification): Does the explanation correctly identify why the vulnerability exists?
- **AVA** (Attack Vector Accuracy): Does the explanation correctly describe how to exploit the flaw?
- **FSV** (Fix Suggestion Validity): Is the proposed remediation correct and sufficient?

D.3 Human Validation

Twenty responses spanning all transformations and difficulty levels underwent independent review by two security experts. Validators assessed:

- Verdict correctness (binary)
- Target finding accuracy (binary)
- Reasoning quality scores (0-1 scale for RCIR, AVA, FSV)

Inter-rater reliability: verdict $\kappa=0.91$, type match $\kappa=0.84$, reasoning $\kappa=0.78$. Human-judge correlation: Pearson's $\rho=0.87$ ($p<0.001$) with 85% decision agreement.

E SUI Sensitivity Analysis

To assess the robustness of SUI rankings to weight choice, we evaluate model performance under five configurations representing different deployment priorities (Table 3). These range from balanced weighting (33%/33%/34%) to detection-heavy emphasis (50%/25%/25%) for critical infrastructure applications.

Table 4 shows complete SUI scores and rankings under each configuration. Rankings exhibit

Table 3: SUI weight configurations for different deployment priorities.

Config	TDR	Rsn	Prec	Rationale
Balanced	0.33	0.33	0.34	Equal weights
Detection (Default)	0.40	0.30	0.30	Practitioner
Quality-First	0.30	0.40	0.30	Research
Precision-First	0.30	0.30	0.40	Production
Detection-Heavy	0.50	0.25	0.25	Critical infra

high stability: average Spearman's $\rho = 0.949 \pm 0.047$ across all configuration pairs (range: [0.829, 1.000]). Grok 4 consistently ranks first across all five configurations. The top-2 positions remain unchanged (Grok 4, Gemini 3 Pro) except in Quality-First weighting, where Claude Opus 4.5's perfect reasoning scores (RCIR/AVA/FSV = 1.0) elevate it to second place.

This high correlation ($\rho > 0.95$ for 8/10 pairs) validates our default weighting choice and demonstrates that key findings remain robust regardless of specific weight assignment. The lowest correlation (0.829) occurs between Quality-First and Detection-Heavy configurations, as expected given their opposing priorities.

F Metric Definitions and Mathematical Framework

F.1 Notation

F.2 Classification Metrics

Standard binary classification metrics:

$$\text{Accuracy} = \frac{TP + TN}{N} \quad (1)$$

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN} \quad (2)$$

$$F_1 = \frac{2 \cdot \text{Prec} \cdot \text{Rec}}{\text{Prec} + \text{Rec}}, \quad F_2 = \frac{5 \cdot \text{Prec} \cdot \text{Rec}}{4 \cdot \text{Prec} + \text{Rec}} \quad (3)$$

where TP , TN , FP , FN denote true/false positives/negatives.

F.3 Target Detection Metrics

Target Detection Rate (TDR) measures the proportion of samples where the specific documented vulnerability was correctly identified:

$$\text{TDR} = \frac{|\{i \in \mathcal{D} \mid \text{target_found}_i = \text{True}\}|}{|\mathcal{D}|} \quad (4)$$

Table 4: Model SUI scores and rankings (in parentheses) under different weight configurations.

Model	Balanced	Default	Quality-First	Precision-First	Detection-Heavy
Grok 4	0.643 (1)	0.625 (1)	0.679 (1)	0.631 (1)	0.596 (1)
Gemini 3 Pro	0.458 (2)	0.448 (2)	0.496 (3)	0.438 (2)	0.429 (2)
Claude Opus 4.5	0.445 (3)	0.423 (3)	0.503 (2)	0.418 (3)	0.386 (4)
GPT-5.2	0.428 (4)	0.414 (4)	0.468 (4)	0.408 (4)	0.389 (3)
Llama 3.1 405B	0.359 (5)	0.333 (5)	0.426 (5)	0.328 (5)	0.290 (5)
DeepSeek v3.2	0.264 (6)	0.253 (6)	0.302 (6)	0.244 (6)	0.233 (6)

Table 5: Core notation for evaluation metrics.

Symbol	Definition
\mathcal{D}	Dataset of all samples
N	Total number of samples ($ \mathcal{D} $)
c_i	Contract code for sample i
v_i	Ground truth vulnerability type for sample i
\mathcal{M}	Model/detector being evaluated
r_i	Model response for sample i
\hat{y}_i	Predicted verdict (vulnerable/safe) for sample i
y_i	Ground truth verdict for sample i
\mathcal{F}_i	Set of findings reported for sample i
$\mathcal{F}_i^{\text{correct}}$	Subset of correct findings for sample i
$\mathcal{F}_i^{\text{hallucinated}}$	Subset of hallucinated findings for sample i

A finding is classified as target found if and only if:

- Type match is at least “partial” (vulnerability type correctly identified)
- Location match is at least “partial” (vulnerable function/line correctly identified)

Lucky Guess Rate (LGR) measures the proportion of correct verdicts where the target vulnerability was not actually found:

$$\text{LGR} = \frac{|\{i \mid \hat{y}_i = y_i \wedge \text{target_found}_i = \text{False}\}|}{|\{i \mid \hat{y}_i = y_i\}|} \quad (5)$$

High LGR indicates the model correctly predicts vulnerable/safe status without genuine understanding of the specific vulnerability.

F.4 Finding Quality Metrics

Finding Precision measures the proportion of reported findings that are correct:

$$\text{Finding Precision} = \frac{\sum_{i \in \mathcal{D}} |\mathcal{F}_i^{\text{correct}}|}{\sum_{i \in \mathcal{D}} |\mathcal{F}_i|} \quad (6)$$

Hallucination Rate measures the proportion of completely fabricated findings:

$$\text{Hallucination Rate} = \frac{\sum_{i \in \mathcal{D}} |\mathcal{F}_i^{\text{hallucinated}}|}{\sum_{i \in \mathcal{D}} |\mathcal{F}_i|} \quad (7)$$

F.5 Reasoning Quality Metrics

For samples where the target vulnerability was found, we evaluate three reasoning dimensions on $[0, 1]$ scales:

- **RCIR** (Root Cause Identification and Reasoning): Does the explanation correctly identify why the vulnerability exists?
- **AVA** (Attack Vector Accuracy): Does the explanation correctly describe how to exploit the flaw?
- **FSV** (Fix Suggestion Validity): Is the proposed remediation correct?

Mean reasoning quality:

$$\bar{R} = \frac{1}{|\mathcal{D}_{\text{found}}|} \sum_{i \in \mathcal{D}_{\text{found}}} \frac{\text{RCIR}_i + \text{AVA}_i + \text{FSV}_i}{3} \quad (8)$$

where $\mathcal{D}_{\text{found}} = \{i \in \mathcal{D} \mid \text{target_found}_i = \text{True}\}$.

F.6 Security Understanding Index (SUI)

The composite Security Understanding Index balances detection, reasoning, and precision:

$$\text{SUI} = w_{\text{TDR}} \cdot \text{TDR} + w_R \cdot \bar{R} + w_{\text{FP}} \cdot \text{Finding Precision} \quad (9)$$

with default weights $w_{\text{TDR}} = 0.40$, $w_R = 0.30$, $w_{\text{FP}} = 0.30$.

Rationale for Weights:

- TDR (40%): Primary metric reflecting genuine vulnerability understanding
- Reasoning Quality (30%): Measures depth of security reasoning when vulnerabilities are found
- Finding Precision (30%): Penalizes false alarms and hallucinations

F.7 Statistical Validation

Ranking Stability. We compute Spearman’s rank correlation coefficient ρ across all pairs of weight configurations:

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)} \quad (10)$$

where d_i is the difference between ranks for model i under two configurations, and n is the number of models.

Human Validation. Inter-rater reliability measured using Cohen's kappa:

$$\kappa = \frac{p_o - p_e}{1 - p_e} \quad (11)$$

where p_o is observed agreement and p_e is expected agreement by chance.

Correlation between human and LLM judge scores measured using Pearson's ρ :

$$\rho = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}} \quad (12)$$