

# Do Frontier LLMs Truly Understand Smart Contract Vulnerabilities?

Anonymous ACL submission

## Abstract

Frontier large language models achieve state-of-the-art performance on code understanding benchmarks, yet their capacity for smart contract security remains unclear. Can they genuinely reason about vulnerabilities, or merely pattern-match against memorized exploits? We introduce **BlockBench**, a benchmark designed to answer this question, revealing heterogeneous capabilities. While some models demonstrate robust semantic understanding, most exhibit substantial surface pattern dependence.

## 1 Introduction

Smart contract vulnerabilities represent one of the most costly security challenges in modern computing. As shown in Figure 1, cryptocurrency theft has resulted in over \$14 billion in losses since 2020, with 2025 already reaching \$3.4 billion, the highest since the 2022 peak (Chainalysis, 2025). The Bybit breach alone accounted for \$1.5 billion, while the Cetus protocol lost \$223 million in minutes due to a single overflow vulnerability (Tsentsura, 2025).




Figure 1: Annual cryptocurrency theft losses (2020–2025). Data from Chainalysis.

Meanwhile, large language models have achieved remarkable success on programming tasks. Frontier models now pass technical interviews, generate production code, and identify bugs across diverse codebases. This raises a natural question: *can these models apply similar expertise to blockchain security?* And if they can, *are they genuinely reasoning about vulnerabilities, or merely pattern-matching against memorized examples?*

This distinction matters. A model that has memoized the 2016 DAO reentrancy attack may flag similar patterns, yet fail when the same flaw appears in unfamiliar syntax. We introduce **BlockBench**, a benchmark designed to answer this question. Our contributions include:

1. **BlockBench**, comprising 263 Solidity vulnerability samples with systematic contamination control and gold standard examples from recent professional security audits.
2. **Composite evaluation metrics** distinguishing genuine understanding from memorization, validated through multi-configuration sensitivity analysis (Spearman’s  $\rho=1.000$ ).
3. **Systematic assessment** revealing 58% best-case detection on mixed samples collapsing to 20% on uncontaminated professional audits, exposing heterogeneous robustness and accuracy-understanding gaps across models.

## 2 Related Work

### 2.1 Traditional Analysis Tools

Early approaches to smart contract vulnerability detection relied on static analysis and symbolic execution. Tools such as Slither (Feist et al., 2019), Mythril (Mueller, 2017), and Security (Tsankov et al., 2018) demonstrated strong precision on syntactically well-defined vulnerability classes. Durieux et al. (2020) conducted a comprehensive evaluation of nine such tools across 47,587 Ethereum contracts, revealing consistent performance on reentrancy and integer overflow detection, yet persistent struggles with vulnerabilities requiring semantic reasoning about contract logic. Ghaleb and Pattabiraman (2020) corroborated these findings, observing that rule-based approaches fundamentally cannot capture the contextual nuances that distinguish exploitable flaws from benign code patterns.

## 070 2.2 LLM-Based Approaches

071 Large language models introduced new possibilities  
072 for bridging this semantic gap. Initial investigations  
073 by Chen et al. (2023) explored prompting  
074 strategies for vulnerability detection, achieving de-  
075 tection rates near 40% while noting pronounced  
076 sensitivity to superficial features such as variable  
077 naming conventions. GPTScan (Sun et al., 2024b)  
078 combined GPT-4 with program analysis to achieve  
079 78% precision on logic vulnerabilities, leveraging  
080 static analysis to validate LLM-generated can-  
081 didates. Sun et al. (2024a) introduced retrieval-  
082 augmented approaches that provide models with  
083 relevant vulnerability descriptions, substantially  
084 improving detection performance. Multi-agent  
085 architectures emerged as another direction, with  
086 systems like GPTLens (Hu et al., 2023) employ-  
087 ing auditor-critic pairs to enhance analytical con-  
088 sistency. Fine-tuning on domain-specific corpora  
089 has yielded incremental gains, though performance  
090 characteristically plateaus below the 85% threshold  
091 regardless of training scale.

## 092 2.3 Pattern Recognition Versus 093 Understanding

094 Beneath these encouraging metrics lies a more  
095 fundamental question: whether observed improve-  
096 ments reflect genuine comprehension of vulnerabil-  
097 ity mechanics or increasingly sophisticated pattern  
098 recognition. Several empirical observations sug-  
099 gest the latter warrants serious consideration. Sun  
100 et al. (2024a) demonstrated that decoupling vulner-  
101 ability descriptions from code context precipitates  
102 catastrophic performance degradation, indicating  
103 that models may rely on memorized associations  
104 between textual cues and vulnerability labels rather  
105 than reasoning about exploit mechanics. Hu et al.  
106 (2023) observed that models produce divergent out-  
107 puts for identical queries even at temperature zero,  
108 a phenomenon difficult to reconcile with deter-  
109 ministic security reasoning. Wu et al. (2024) showed  
110 through counterfactual tasks in adjacent domains  
111 that language models systematically fail when fa-  
112 miliar patterns are disrupted, defaulting to memo-  
113 rized responses rather than applying causal logic to  
114 novel configurations.

## 115 2.4 Evaluation Methodology

116 The distinction between pattern recognition and  
117 genuine understanding carries profound implica-  
118 tions for security applications, where adversarial

actors actively craft exploits to evade detection.  
A model that has memorized the surface features  
of known vulnerabilities provides little defense  
against novel attack vectors or obfuscated variants  
of familiar exploits. Existing benchmarks such  
as SmartBugs Curated (Durieux et al., 2020) and  
DeFiVulnLabs (SunWeb3Sec, 2023) assess binary  
detection outcomes without examining whether  
models can identify specific code elements that  
enable exploitation, distinguish genuine vulnerabil-  
ties from superficially suspicious but benign pat-  
terns, or maintain accuracy when surface-level cues  
are systematically removed. Our work contributes  
evaluation methodology that directly probes this  
distinction through adversarial transformations pre-  
serving vulnerability semantics while removing sur-  
face cues.

## 3 BlockBench

We introduce BlockBench, a benchmark for eval-  
uating whether AI models genuinely understand  
smart contract vulnerabilities. The benchmark is  
designed to distinguish genuine security under-  
standing from pattern memorization, comprising  
290 vulnerable Solidity contracts with 322 trans-  
formation variants, spanning over 30 vulnerability  
categories (Appendix B).

Let  $\mathcal{D}$  represent the dataset, where  $\mathcal{D} = \{(c_i, v_i, m_i)\}_{i=1}^{290}$ . Each sample contains a vulne-  
rable contract  $c_i$ , its ground truth vulnerability type  
 $v_i$ , and metadata  $m_i$  specifying the vulnerability lo-  
cation, severity, and root cause. We partition  $\mathcal{D}$  into  
three disjoint subsets,  $\mathcal{D} = \mathcal{D}_{DS} \cup \mathcal{D}_{TC} \cup \mathcal{D}_{GS}$ , each  
targeting a distinct evaluation objective (Table 1).

Subset	N	Sources
Difficulty Stratified (DS)	210	SmartBugs, DeFiVulnLabs
Temporal Contamination (TC)	46	Real-world exploits
Gold Standard (GS)	34	Code4rena, Spearbit

Table 1: BlockBench composition by subset and primary sources.

**Difficulty Stratified.**  $\mathcal{D}_{DS}$  draws from estab-  
lished vulnerability repositories including Smart-  
Bugs Curated (Ferreira et al., 2020), Trail of Bits’  
Not So Smart Contracts (Trail of Bits, 2018), and  
DeFiVulnLabs (SunWeb3Sec, 2023). Samples are  
stratified into four difficulty tiers based on detection  
complexity, with distribution {86, 81, 30, 13} from  
Tier 1 (basic patterns) through Tier 4 (expert-level  
vulnerabilities requiring deep protocol knowledge).

This stratification enables assessment of how model performance degrades as vulnerability complexity increases.

**Temporal Contamination.**  $\mathcal{D}_{TC}$  reconstructs 46 real-world DeFi exploits spanning 2016 to 2024, representing over \$1.65 billion in documented losses. Notable incidents include The DAO (\$60M, 2016), Nomad Bridge (\$190M, 2022), and Curve Vyper (\$70M, 2023). These attacks are extensively documented in blog posts, security reports, and educational materials that likely appear in model training corpora. To probe whether models genuinely understand these vulnerabilities or merely recognize them, we apply systematic transformations that preserve vulnerability semantics while removing surface cues (detailed in §4).

**Gold Standard.**  $\mathcal{D}_{GS}$  derives from 34 professional security audit findings by Code4rena (Code4rena, 2025), Spearbit (Spearbit, 2025), and MixBytes (MixBytes, 2025) disclosed after September 2025. We designate this subset as “gold standard” because all samples postdate  $t_{cutoff} = \text{August 2025}$ , the most recent training cut-off among frontier models evaluated in this work. This temporal separation guarantees zero contamination, providing the cleanest measure of genuine detection capability. The subset emphasizes logic errors (53%) and includes 10 high-severity and 24 medium-severity findings.

These complementary subsets collectively enable rigorous assessment of both detection capability and the distinction between pattern memorization and genuine security understanding.

## 4 Methodology

Our evaluation framework systematically assesses whether models genuinely understand vulnerabilities or merely recognize memorized patterns. Figure 2 illustrates the complete pipeline.

### 4.1 Adversarial Transformations

To distinguish pattern memorization from genuine understanding, we apply semantic-preserving transformations to  $\mathcal{D}_{TC}$ . Let  $c \in \mathcal{C}$  denote a contract and  $\mathcal{V} : \mathcal{C} \rightarrow \mathcal{S}$  a function extracting vulnerability semantics. A transformation  $\mathcal{T} : \mathcal{C} \rightarrow \mathcal{C}$  is *semantic-preserving* iff  $\mathcal{V}(\mathcal{T}(c)) = \mathcal{V}(c)$ . We define eight transformations targeting distinct recognition pathways, organized hierarchically in Figure 3.

❖ **Sanitization** ( $\mathcal{T}_S$ ). Removes protocol-identifying information through 280+ pattern replacements:  $\mathcal{T}_S(c) = \text{replace}(c, \mathcal{P}_{\text{protocol}}, \mathcal{P}_{\text{generic}})$  where  $\mathcal{P}_{\text{protocol}}$  maps protocol-specific identifiers (e.g., NomadReplica) to generic equivalents (e.g., BridgeReplica). Tests whether detection relies on recognizing known protocol names.

☒ **No-Comments** ( $\mathcal{T}_N$ ). Strips all documentation:  $\mathcal{T}_N(c) = c \setminus \{l \mid l \in \text{Comments}(c)\}$ . Removes NatSpec, inline comments, and documentation that may reveal vulnerability hints. Tests pure code analysis capability.

⊕ **Chameleon** ( $\mathcal{T}_C$ ). Applies domain-shifting vocabulary while preserving logic:  $\mathcal{T}_C(c) = \text{replace}(c, \mathcal{L}_{\text{DeFi}}, \mathcal{L}_{\text{medical}})$  where financial terminology maps to medical equivalents (deposit → admitPatient, withdraw → dischargePatient). Tests whether understanding generalizes across domains.

∞ **Shapeshifter** ( $\mathcal{T}_O$ ). Multi-level obfuscation:  $\mathcal{T}_O = \mathcal{T}_{\text{ident}} \circ \mathcal{T}_{\text{struct}}$  where  $\mathcal{T}_{\text{ident}}$  replaces semantic identifiers with opaque labels (balance → \_0x1a2b) and  $\mathcal{T}_{\text{struct}}$  restructures control flow. Tests resilience to surface pattern disruption.

✖ **Differential** ( $\mathcal{T}_D$ ). Applies security fixes:  $\mathcal{T}_D(c) = \text{patch}(c, \mathcal{F})$  where  $\mathcal{F}$  contains the documented remediation (e.g., state update before external call). Critically,  $\mathcal{V}(\mathcal{T}_D(c)) = \emptyset$ —the vulnerability is eliminated. Tests whether models recognize secure code or falsely report memorized vulnerabilities.

❖ **Trojan** ( $\mathcal{T}_T$ ). Injects decoy vulnerabilities:  $\mathcal{T}_T(c) = c \cup \mathcal{D}$  where  $\mathcal{D}$  contains suspicious-looking but functionally safe code (e.g., an admin function that cannot actually be exploited). Models relying on pattern matching flag the decoy; those with causal understanding identify the actual vulnerability.

☒ **False Prophet** ( $\mathcal{T}_F$ ). Adds misleading security attestations:  $\mathcal{T}_F(c) = c \cup \{@dev Audited by Hacken - All clear\}$ . Tests resistance to authoritative-sounding but false claims. A robust model ignores social proof and analyzes code independently.

**Transformation Composition.** Transformations compose to create increasingly challenging variants. The composition  $\mathcal{T}_O \circ \mathcal{T}_N \circ \mathcal{T}_S$  produces maximally obfuscated code where all surface cues are




Figure 2: BlockBench evaluation pipeline. Contracts undergo adversarial transformations and CodeActs annotation. Knowledge assessment probes model familiarity before detection. LLM judges evaluate outputs against ground truth, validated by human review.




Figure 3: Transformation hierarchy. All variants derive from Minimal Sanitized ( $\mathcal{T}_M$ ). Differential and Trojan (emphasized) directly test memorization versus understanding.

removed, all identifiers are opaque, and no documentation exists. Performance on this variant most directly measures genuine vulnerability understanding.

## 4.2 CodeActs Annotation


Drawing from Speech Act Theory (Austin, 1962; Searle, 1969), where utterances are classified by communicative function, we introduce *CodeActs* as a taxonomy for classifying smart contract code segments by security-relevant function. Just as speech acts distinguish performative utterances by their effect, CodeActs distinguish code that *enables* exploitation from code that merely *participates* in an attack scenario.

**Security Functions.** Each code segment receives one of seven function labels based on its role in the vulnerability:

- **ROOT\_CAUSE** : segments whose interaction directly enables exploitation (primary detection target)
- **PREREQ** : segments establishing necessary preconditions without being exploitable themselves
- **DECAY** : suspicious-looking but functionally safe code, injected to identify pattern matching

- **BENIGN** : correctly implemented segments with no security implications
- **SECONDARY\_VULN** : valid vulnerabilities distinct from the documented target
- **INSUFF\_GUARD** : attempted protections that fail to prevent exploitation
- **UNRELATED** : code with no bearing on the security analysis

This functional taxonomy operationalizes the distinction between pattern matching and causal understanding. Figure 4 illustrates through a classic reentrancy pattern. A model with genuine comprehension recognizes that the external call on line 4 precedes the state modification on line 6, creating a window for recursive exploitation. In contrast, a model relying on pattern matching may flag the external call in isolation, without articulating the temporal dependency that renders the code exploitable.



Legend: PREREQ ROOT\_CAUSE BENIGN

Figure 4: CodeActs annotation for reentrancy. Lines 4 and 6 (**ROOT\_CAUSE**) enable exploitation through their ordering; line 2 (**PREREQ**) establishes preconditions.

A correct detection must identify **ROOT\_CAUSE** segments and explain their causal relationship. Flagging only line 4, or failing to articulate why the ordering matters, reveals incomplete understanding despite a nominally correct vulnerability classification.

**Annotation Variants.** CodeActs enable three evaluation strategies targeting different aspects of model comprehension:

- **Minimal Sanitized** ( $\mathcal{T}_M$ ) establishes baseline detection with **ROOT\_CAUSE** and

- 311            PREREQ annotations only  
 312    • **Trojan** ( $\mathcal{T}_T$ ) injects DECOY segments that  
 313      appear vulnerable but lack exploitability  
 314    • **Differential** ( $\mathcal{T}_D$ ) presents fixed code where  
 315      former ROOT\_CAUSE becomes BENIGN

316    Models that flag DECOY segments reveal  
 317      pattern-matching behavior. Models that report  
 318      vulnerabilities in Differential variants, where the  
 319      fix converts ROOT\_CAUSE to BENIGN, demon-  
 320      strate memorization of the original exploit rather  
 321      than analysis of the presented code.

322    We define 17 security-relevant code operations  
 323      (e.g., EXT\_CALL, STATE\_MOD, ACCESS\_CTRL),  
 324      each receiving a security function label based on  
 325      its role. The same operation type can have different  
 326      functions depending on context: an EXT\_CALL  
 327      might be ROOT\_CAUSE in reentrancy, PREREQ  
 328      in oracle manipulation, or DECOY when delib-  
 329      erately injected. The full taxonomy appears in  
 330      Appendix C.

### 331    4.3 Detection Protocol

332    We evaluate seven frontier models spanning seven  
 333      AI labs: Claude Opus 4.5 (Anthropic), GPT-5.2  
 334      (OpenAI), Gemini 3 Pro (Google), DeepSeek v3.2  
 335      (DeepSeek), Llama 4 Maverick (Meta), Grok 4  
 336      Fast (xAI), and Qwen3-Coder-Plus (Alibaba). This  
 337      selection ensures one flagship representation per  
 338      major AI lab, covering both general-purpose mod-  
 339      els and a code-specialized variant.

340    For DS and TC datasets, models receive a direct  
 341      zero-shot prompt requesting structured JSON out-  
 342      put with vulnerability type, location, root cause,  
 343      attack scenario, and fix. For GS, we additionally  
 344      test five prompting strategies: *zero-shot* (baseline),  
 345      *context-enhanced* (with brief protocol documenta-  
 346      tion), *chain-of-thought* (explicit step-by-step rea-  
 347      soning), *naturalistic* (informal code review), and  
 348      *adversarial* (misleading priming suggesting prior  
 349      audit approval). All evaluations use temperature 0.  
 350      Detailed prompt descriptions and templates appear  
 351      in Appendix F.

### 352    4.4 Knowledge Assessment

353    Before detection, we probe whether models possess  
 354      prior knowledge of documented exploits by query-  
 355      ing for factual details (date, amount lost, vulnerabil-  
 356      ity type, attack mechanism). Since models may hal-  
 357      lucinate familiarity, we validate responses against  
 358      ground truth metadata. Let  $\mathcal{K}(m, e) \in \{0, 1\}$  indi-  
 359      cate *verified* knowledge, requiring accurate recall

of at least two factual details. This enables diag-  
 360      nóstic interpretation:  $\mathcal{K} = 1$  with detection failure  
 361      under obfuscation ( $\mathcal{T}_O$ ) indicates memorization;  
 362       $\mathcal{K} = 1$  with robust detection across transforma-  
 363      tions indicates understanding;  $\mathcal{K} = 0$  with successful de-  
 364      tection indicates genuine analytical capability.

## 365    4.5 LLM-as-Judge Evaluation

366    LLM judges evaluate detection outputs against  
 367      ground truth. A finding qualifies as TARGET\_MATCH if it correctly identifies the root cause  
 368      mechanism, vulnerable location, and type classi-  
 369      fication; PARTIAL\_MATCH for correct root cause  
 370      with imprecise type; BONUS\_VALID for valid find-  
 371      ings beyond documented ground truth. Invalid find-  
 372      ings are classified as HALLUCINATED, MISCHAR-  
 373      ACTERIZED, DESIGN\_CHOICE, OUT\_OF\_SCOPE,  
 374      SECURITY\_THEATER, or INFORMATIONAL.

375    For matched findings, judges assess explana-  
 376      tion quality on three dimensions (0-1 scale): *Root*  
 377      *Cause Identification Rate* (RCIR) measures articu-  
 378      lation of the exploitation mechanism; *Attack Vector*  
 379      *Validity* (AVA) assesses whether attack scenarios  
 380      are concrete and executable; *Fix Suggestion Valid-  
 381      ity* (FSV) evaluates remediation effectiveness.

382    Three judge models (GLM-4.7, Mistral Large,  
 383      MIMO v2) independently evaluate each output,  
 384      reducing individual bias and enabling inter-judge  
 385      agreement measurement. A subset undergoes ex-  
 386      pert review to calibrate automated judgment, with  
 387      reliability measured using Cohen’s  $\kappa$  for classifi-  
 388      cation and Spearman’s  $\rho$  for quality scores (Ap-  
 389      pendix G).

## 390    4.6 Evaluation Metrics

391    **Target Detection Rate (TDR).** Primary metric:  
 392       $TDR = |\{s : \text{TARGET\_MATCH}(s)\}| / |\mathcal{D}|$ . Mea-  
 393      sures correct identification of documented vulnera-  
 394      bilities with matching root cause and location.


395    **Quality Metrics.** For detected targets, we report  
 396      mean RCIR, AVA, and FSV. These distinguish  
 397      shallow pattern matches from deep understanding  
 398      through accurate root cause analysis, concrete at-  
 399      tack scenarios, and valid remediations.

400    **Transformation Degradation.** We compute  
 401       $\Delta_{\mathcal{T}} = TDR(c) - TDR(\mathcal{T}(c))$  for each transfor-  
 402      mation. Significant degradation despite claimed  
 403      knowledge ( $\mathcal{K} = 1$ ) provides evidence for mem-  
 404      orization. We apply McNemar’s test for paired  
 405      comparisons and report effect sizes.

## 408 5 Results

409 We evaluate six frontier models on  $n = 58$  samples  
 410 from BlockBench: 10 GS, 20 TC, 28 DS.

### 411 5.1 Overall Performance



412 Figure 5: Target Detection Rate across all models. Best  
 413 performer achieves 58% detection, while highest accu-  
 414 racy (88%) corresponds to lowest TDR (18%).

415 Table 2 and Figure 5 show aggregate performance  
 416 by TDR. Gemini 3 Pro achieves highest  
 417 detection (58%), followed by GPT-5.2 (56%) and  
 418 Claude Opus 4.5 (53%). Combining detection, rea-  
 419 soning, and precision into SUI, GPT-5.2 ranks first  
 420 (0.746) on finding precision (77%).

421 Llama 3.1 405B exhibits severe accuracy-TDR  
 422 gap: 88% accuracy yet 18% TDR, classifying  
 423 samples as vulnerable without identifying specific  
 424 flaws. This 70pp discrepancy shows binary classifi-  
 425 cation inadequately measures security under-  
 426 standing. Models achieving target detection show strong  
 427 reasoning (RCIR/AVA/FSV  $\geq 0.95$ ).

### 428 5.2 Gold Standard Performance

429 Gold Standard samples from post-September 2025  
 430 audits ensure zero temporal contamination. Perfor-  
 431 mance drops substantially: Claude Opus 4.5 leads  
 432 (20% TDR), followed by Gemini 3 Pro (11%),  
 433 GPT-5.2 (10%), Grok 4 (10%). DeepSeek v3.2  
 434 and Llama detect zero targets. Models experience  
 435 34-50pp drops from overall to Gold Standard.

### 436 5.3 Transformation Robustness

437 **Sanitization.** Neutralizing security-suggestive  
 438 identifiers causes variable degradation. GPT-5.2  
 439 and DeepSeek v3.2 maintain performance, while  
 440 Grok 4 drops 40pp, exposing varying lexical  
 441 reliance.


442 **Domain Shift.** Replacing blockchain terminol-  
 443 ogy with medical vocabulary shows mixed impact  
 444 (20-60% TDR). GPT-5.2 maintains 60% detection  
 445 while others degrade 20-50%. 470

**Prompt Framing.** Performance varies across  
 direct, adversarial, and naturalistic prompts. Gemini 3 Pro and GPT-5.2 show robustness (18-21pp  
 drops), while Claude Opus 4.5 and DeepSeek v3.2  
 degrade more (21-39pp). Llama exhibits inconsis-  
 tent behavior (adversarial: 0%, naturalistic: 25%).

### 446 5.4 Human Validation

447 Two security experts validated 31 samples (116  
 448 comparisons across models) with 92% agreement  
 $(\kappa=0.84, \rho=0.85, p<0.0001)$ . Judge achieved per-  
 449 fect recall with 84% precision ( $F1=0.91$ ).

## 450 6 Discussion



451 Figure 6: Security Understanding Index trajectory  
 452 across progressive transformations of TC samples. GPT-  
 453 5.2 maintains near-constant performance while most  
 454 models degrade, revealing varying degrees of surface  
 455 pattern reliance.

**Understanding versus Memorization.** Figure 6 reveals heterogeneous robustness across models. GPT-5.2 maintains stable SUI ( $78.0 \rightarrow 77.2$ ) through sanitization, domain shifts, and obfuscation, demonstrating genuine semantic understanding. In contrast, DeepSeek v3.2 degrades 19.7 points ( $78.7 \rightarrow 59.0$ ), indicating surface pattern dependence. Most models exhibit intermediate behavior, leveraging lexical cues when available while retaining partial structural understanding (Chen et al., 2021; Wu et al., 2024). This heterogeneity suggests current training methods produce inconsistent abstraction capabilities across architectures (Sánchez Salido et al., 2025). While genuine security understanding is demonstrably possible, most frontier models have not achieved it.

Model	TDR	SUI	Acc	RCIR	AVA	FSV	Findings
Gemini 3 Pro	<b>57.6</b>	0.734	<b>93.9</b>	0.97	0.97	0.95	2.6
GPT-5.2	55.9	<b>0.746</b>	75.0	0.97	0.98	<b>0.97</b>	2.4
Claude Opus 4.5	52.9	0.703	83.8	0.98	0.99	0.97	3.5
Grok 4	44.1	0.677	69.1	<b>0.98</b>	<b>1.00</b>	<b>0.97</b>	2.1
DeepSeek v3.2	38.2	0.599	82.4	0.91	0.92	0.86	3.0
Llama 3.1 405B	17.9	0.393	<b>88.1</b>	0.88	0.90	0.83	2.0

Table 2: Overall performance ranked by Target Detection Rate. Best values bold.

471  
472 **Measurement Inadequacy.** The accuracy-TDR  
473 gap exposes fundamental metric limitations. Llama  
474 3.1 405B achieves 88% accuracy yet only 18%  
475 TDR, correctly classifying samples as vulnerable  
476 without identifying specific flaw types or locations  
477 ([Jimenez et al., 2024](#)). For security practitioners  
478 requiring actionable findings, binary classification  
479 provides insufficient value. Effective evaluation  
480 must measure precise vulnerability localization, not  
merely anomaly detection.

481 **Practical Implications.** Current frontier models  
482 cannot serve as autonomous auditors. Best per-  
483 formance reaches 58% detection with substantial  
484 Gold Standard degradation (20% maximum). How-  
485 ever, complementary strengths suggest ensemble  
486 potential: Grok 4 offers breadth, GPT-5.2 provides  
487 consistency, Claude delivers explanation quality.  
488 Workflows positioning LLMs as assistive tools with  
489 mandatory expert review align capabilities with cur-  
490 rent limitations ([Hu et al., 2023](#); [Ince et al., 2025](#)).

## 491 7 Conclusion

492 BlockBench evaluates whether frontier LLMs gen-  
493 uinely understand smart contract vulnerabilities or  
494 merely pattern-match. Our assessment of six fron-  
495 tier models reveals substantial limitations. Best per-  
496 formance reaches 58% detection on mixed samples,  
497 collapsing to 20% on Gold Standard audits. Llama  
498 3.1 405B achieves 88% accuracy yet 18% TDR,  
499 demonstrating binary classification inadequately  
500 measures security understanding.

501 Models exhibit heterogeneous robustness. While  
502 GPT-5.2 maintains stable performance across trans-  
503 formations, most models degrade when surface  
504 cues are removed. Current frontier LLMs cannot  
505 serve as autonomous auditors but show promise in  
506 ensemble workflows with mandatory expert review.  
507 Future work should develop sanitization-resistant  
508 methods and explore hybrid LLM-verification ar-  
chitectures.

## 510 Ethical Considerations

511 BlockBench poses dual-use risks: adversarial  
512 prompts demonstrate methods that could suppress  
513 detection, while detailed vulnerability documen-  
514 tation may assist malicious actors. We justify public  
515 release on several grounds: adversarial robustness  
516 represents a fundamental requirement for security  
517 tools, malicious actors will discover these vulner-  
518 abilities regardless, and responsible disclosure en-  
519 ables proactive mitigation. All samples derive from  
520 already-disclosed vulnerabilities and public secu-  
521 rity audits, ensuring no novel exploit information is  
522 revealed. Practitioners should avoid over-reliance  
523 on imperfect tools, as false negatives create secu-  
524 rity gaps while false confidence may reduce manual  
525 review rigor.

## 526 Limitations and Future Work

527 Our evaluation uses 58 samples, including 10 Gold  
528 Standard examples from recent professional au-  
529 dit. We assess zero-shot prompting exclusively  
530 and provide models only with the contract code  
531 necessary to expose each vulnerability. In real audit  
532 settings, analysts often rely on additional semantic  
533 context such as protocol goals, intended invariants,  
534 expected economic behavior, and threat models.  
535 Providing this context may improve vulnerability  
536 detection, particularly for logic-related flaws in the  
537 Gold Standard subset.

538 Future work should explore chain-of-thought  
539 reasoning, retrieval-augmented analysis, and ex-  
540 plicit specification of protocol intent to better cap-  
541 ture contextual information. It should also expand  
542 sample diversity across blockchain ecosystems, de-  
543 velop sanitization-resistant analysis using control-  
544 flow and data-flow representations, and explore hy-  
545 brid LLM-verification architectures that integrate  
546 formal specifications and contextual reasoning ([Liu  
et al., 2024](#)).

## AI Assistance

Claude Sonnet 4.5 assisted with evaluation pipeline code and manuscript refinement. All research design, experimentation, and analysis were conducted by the authors.

## References

- J. L. Austin. 1962. *How to Do Things with Words*. Oxford University Press.
- Chainalysis. 2025. Crypto theft reaches \$3.4b in 2025. <https://www.chainalysis.com/blog/crypto-hacking-stolen-funds-2026/>. Accessed: 2025-12-18.
- Chong Chen, Jianzhong Nie, Xingyu Peng, Jian Yang, Dan Wang, Jiayuan Zhuo, Zhenqi Liu, and Zhun Yang. 2023. When ChatGPT meets smart contract vulnerability detection: How far are we? *arXiv preprint arXiv:2309.05520*.
- Mark Chen et al. 2021. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.
- Code4rena. 2025. Competitive audit contest findings. <https://code4rena.com>.
- Thomas Durieux, João F. Ferreira, Rui Abreu, and Pedro Cruz. 2020. Empirical review of automated analysis tools on 47,587 Ethereum smart contracts. In *Proceedings of the ACM/IEEE 42nd International Conference on Software Engineering*, pages 530–541.
- Josselin Feist, Gustavo Grieco, and Alex Groce. 2019. Slither: A static analysis framework for smart contracts. In *Proceedings of the 2nd International Workshop on Emerging Trends in Software Engineering for Blockchain*, pages 8–15.
- João F. Ferreira, Pedro Cruz, Thomas Durieux, and Rui Abreu. 2020. Smartbugs: A framework to analyze Solidity smart contracts. In *Proceedings of the 35th IEEE/ACM International Conference on Automated Software Engineering*, pages 1349–1352.
- Asem Ghaleb and Karthik Pattabiraman. 2020. How effective are smart contract analysis tools? Evaluating smart contract static analysis tools using bug injection. In *Proceedings of the 29th ACM SIGSOFT International Symposium on Software Testing and Analysis*, pages 415–427.
- S M Mostaq Hossain et al. 2025. Leveraging large language models and machine learning for smart contract vulnerability detection. *arXiv preprint arXiv:2501.02229*.
- Sihao Hu, Tiansheng Huang, Feiyang Liu, Sunjun Ge, and Ling Liu. 2023. Large language model-powered smart contract vulnerability detection: New perspectives. *arXiv preprint arXiv:2310.01152*.
- Peter Ince, Jiangshan Yu, Joseph K. Liu, Xiaoning Du, and Xiapu Luo. 2025. Gendetect: Generative large language model usage in smart contract vulnerability detection. In *Provable and Practical Security (ProvSec 2025)*. Springer.
- Carlos E. Jimenez et al. 2024. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*.
- Ye Liu, Yue Xue, Daoyuan Wu, Yuqiang Sun, Yi Li, Miaolei Shi, and Yang Liu. 2024. Propertygpt: LLM-driven formal verification of smart contracts through retrieval-augmented property generation. *arXiv preprint arXiv:2405.02580*.
- MixBytes. 2025. Smart contract security audits. <https://mixbytes.io/audit>.
- Bernhard Mueller. 2017. Mythril: Security analysis tool for Ethereum smart contracts. <https://github.com/ConsenSys/mythril>.
- Eva Sánchez Salido, Julio Gonzalo, and Guillermo Marco. 2025. None of the others: a general technique to distinguish reasoning from memorization in multiple-choice llm evaluation benchmarks. *arXiv preprint arXiv:2502.12896*.
- John R. Searle. 1969. *Speech Acts: An Essay in the Philosophy of Language*. Cambridge University Press.
- Spearbit. 2025. Security audit portfolio. <https://github.com/spearbit/portfolio>.
- Yuqiang Sun, Daoyuan Wu, Yue Xue, Han Liu, Wei Ma, Lyuye Zhang, Miaolei Shi, and Yang Liu. 2024a. LLM4Vuln: A unified evaluation framework for decoupling and enhancing LLMs' vulnerability reasoning. *arXiv preprint arXiv:2401.16185*.
- Yuqiang Sun, Daoyuan Wu, Yue Xue, Han Liu, Haijun Wang, Zhengzi Xu, Xiaofei Xie, and Yang Liu. 2024b. When GPT meets program analysis: Towards intelligent detection of smart contract logic vulnerabilities in GPTScan. In *ICSE*.
- SunWeb3Sec. 2023. DeFiVulnLabs: Learn common smart contract vulnerabilities. <https://github.com/SunWeb3Sec/DeFiVulnLabs>.
- Trail of Bits. 2018. Not so smart contracts: Examples of common Ethereum smart contract vulnerabilities. <https://github.com/crytic/not-so-smart-contracts>.
- Petar Tsankov, Andrei Dan, Dana Drachsler-Cohen, Arthur Gervais, Florian Bünzli, and Martin Vechev. 2018. Securify: Practical security analysis of smart contracts. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*, pages 67–82.

649 Kostiantyn Tsentsura. 2025. [Why DEX exploits cost](#)  
650 [\\$3.1b in 2025: Analysis of 12 major hacks](#). Technical  
651 report, Yellow Network.

652 Zhaofeng Wu, Linlu Qiu, Alexis Ross, Ekin Akyürek,  
653 Boyuan Chen, Bailin Wang, Najoung Kim, Jacob An-  
654 dreas, and Yoon Kim. 2024. Reasoning or reciting?  
655 Exploring the capabilities and limitations of language  
656 models through counterfactual tasks. *arXiv preprint*  
657 *arXiv:2307.02477*.

## A Data and Code Availability

To support reproducibility and future research, we will release all benchmark data and evaluation code upon publication, including 290 base contracts with ground truth annotations, all transformation variants, model evaluation scripts, LLM judge implementation, prompt templates, and analysis notebooks.

## B Vulnerability Type Coverage

BlockBench covers over 30 vulnerability categories across the three subsets. Table 3 shows the primary categories and their distribution.

Vulnerability Type	DS	TC	GS
Access Control	22	14	3
Reentrancy	37	7	—
Logic Error	19	2	18
Unchecked Return	48	—	1
Integer/Arithmetic Issues	16	5	—
Oracle Manipulation	4	8	1
Weak Randomness	8	—	—
DOS	9	—	3
Front Running	5	—	2
Signature Issues	4	1	3
Flash Loan	2	—	2
Honeypot	7	—	—
Other Categories	29	9	1
<b>Total</b>	<b>210</b>	<b>46</b>	<b>34</b>

Table 3: Vulnerability type distribution across BlockBench subsets. “Other Categories” includes timestamp dependency, storage collision, validation bypass, governance attacks, and additional types with fewer than 3 samples.

## C CodeActs Taxonomy

Table 4 presents the complete CodeActs taxonomy with all 17 security-relevant code operations.

**Security Function Assignment.** Each CodeAct in a sample is assigned one of six security functions based on its role:

- **Root\_Cause:** Directly enables exploitation (target)
- **Prereq:** Necessary for exploit but not the cause
- **Insuff\_Guard:** Failed protection attempt
- **Decoy:** Looks vulnerable but is safe (tests pattern-matching)
- **Benign:** Correctly implemented, safe
- **Secondary:** Real vulnerability not in ground truth

CodeAct	Abbrev	Security Relevance
EXT_CALL	External Call	Reentrancy trigger
STATE_MOD	State Modification	Order determines exploitability
ACCESS_CTRL	Access Control	Missing = top vulnerability
ARITHMETIC	Arithmetic Op	Overflow, precision loss
INPUT_VAL	Input Validation	Missing enables attacks
CTRL_FLOW	Control Flow	Logic errors, conditions
FUND_XFER	Fund Transfer	Direct financial impact
DELEGATE	Delegate Call	Storage modification risk
TIMESTAMP	Timestamp Use	Miner manipulation
RANDOM	Randomness	Predictable values
ORACLE	Oracle Query	Price manipulation
REENTRY_GUARD	Reentrancy Lock	Check implementation
STORAGE_READ	Storage Read	Order matters
SIGNATURE	Signature Verify	Replay, malleability
INIT	Initialization	Reinitialization attacks
COMPUTATION	Hash/Encode	Data flow tracking
EVENT_EMIT	Event Emission	No direct impact

Table 4: Complete CodeActs taxonomy (17 security-relevant types).

**Annotation Format.** Each TC sample includes line-level annotations:

```

1 code_acts:
2   - line: 53
3     code_act: INPUT_VAL
4     security_function: ROOT_CAUSE
5     observation: 'messages[hash] == 0 passes
6                   for any unprocessed hash'
```

## D Related Work (Expanded)

**Traditional Smart Contract Analysis.** Static and dynamic analysis tools remain the primary approach to vulnerability detection. Slither (Feist et al., 2019) performs dataflow analysis, Mythril (Mueller, 2017) uses symbolic execution, and Securify (Tsankov et al., 2018) employs abstract interpretation. Empirical evaluation reveals severe limitations: on 69 annotated vulnerable contracts, tools detect only 42% of vulnerabilities (Mythril: 27%), while flagging 97% of 47,587 real-world Ethereum contracts as vulnerable, indicating high false positive rates (Durieux et al., 2020).

**LLM-Based Vulnerability Detection.** Recent work explores LLMs for smart contract analysis. GPTLens (Hu et al., 2023) employs adversarial auditor-critic interactions, while PropertyGPT (Liu et al., 2024) combines retrieval-augmented generation with formal verification. Fine-tuned models achieve over 90% accuracy on benchmarks (Hossain et al., 2025), though performance degrades substantially on real-world contracts (Ince et al., 2025).

**Benchmark Datasets.** SmartBugs Curated (Ferreira et al., 2020) provides 143 annotated contracts as a standard evaluation dataset, while SolidiFI (Ghaleb and Pattabiraman, 2020) uses bug injec-

tion to create controlled samples. Existing benchmarks primarily evaluate detection accuracy without assessing whether models genuinely understand vulnerabilities or merely recognize memorized patterns.

**LLM Robustness and Memorization.** Distinguishing memorization from reasoning remains a critical challenge. Models exhibit high sensitivity to input modifications, with performance drops of up to 57% on paraphrased questions (Sánchez Salido et al., 2025). Wu et al. (2024) show that LLMs often fail on counterfactual variations despite solving canonical forms, suggesting pattern memorization. Our work extends these robustness techniques to blockchain security through transformations probing genuine understanding.

## E Transformation Specifications

We apply four adversarial transformations to probe whether models rely on surface cues or genuine semantic understanding. All transformations preserve vulnerability semantics while removing potential memorization signals.

### E.1 Sanitization (sn)

Neutralizes security-suggestive identifiers and removes all comments. Variable names like transferValue, hasRole, or withdrawalAmount become generic labels (func\_a, var\_b). Function names follow similar neutralization. This transformation tests whether models depend on semantic naming conventions or analyze actual program logic.

#### Example:

```
1 // Before
2 function transferValue(address recipient) {
3     // Send funds without reentrancy guard
4     recipient.call.value(balance)("");
5 }
6
7 // After (Sanitized)
8 function func_a(address param_b) {
9     param_b.call.value(var_c)("");
10 }
```

### E.2 No-Comments (nc)

Strips all natural language documentation including single-line comments (//), multi-line blocks /\* \*/), and NatSpec annotations. Preserves all code structure, identifiers, and logic. Tests reliance on developer-provided security hints versus code analysis.

### E.3 Chameleon (ch)

Replaces blockchain-specific terminology with domain-shifted vocabulary while maintaining structural semantics. Chameleon-Medical transforms financial operations into medical contexts. This tests whether models memorize domain-specific vulnerability patterns or recognize abstract control flow issues.

#### Example transformations:

- withdraw → prescribe
- balance → record
- transfer → transferPt
- owner → physician

### E.4 Shapeshifter (ss)

Applies progressive obfuscation at three levels:

**Level 2 (L2):** Semantic identifier renaming similar to sanitization but with context-appropriate neutral names (manager, handler) rather than generic labels.

**Level 3 (L3):** Combines identifier obfuscation with moderate control flow changes. Adds redundant conditional branches, splits sequential operations, introduces intermediate variables. Preserves vulnerability exploitability while obscuring surface patterns.

#### Example (L3):

```
1 // Original vulnerable pattern
2 if (!authorized) revert();
3 recipient.call.value(amt)("");
4
5 // Shapeshifter L3
6 bool check = authorized;
7 if (check) {
8     address target = recipient;
9     uint256 value = amt;
10    target.call.value(value)("");
11 } else {
12     revert();
13 }
```

These transformations generate 1,343 variants from 263 base samples, enabling systematic robustness evaluation across transformation trajectories.

## F Prompt Templates

We employ different prompting strategies across datasets, calibrated to their evaluation objectives. Table ?? summarizes the strategy matrix.

### F.1 Direct Prompt

Used for DS and TC datasets. Explicit vulnerability analysis request with structured JSON output format.

#### System Prompt (excerpt):

Dataset	Strategy	Context	Protocol	CoT <sup>86</sup>	Framing
DS/TC	Direct	–	–	887 888	Expert
GS	Zero-shot	✓	–	889	Expert
GS	Context-enhanced	✓	✓	–	Expert
GS	Chain-of-thought	✓	✓	✓	Expert
GS	Naturalistic	✓	✓	✓	Casual
GS	Adversarial	✓	✓	✓	Biased

Table 5: Prompting strategy matrix. Context includes related contract files; Protocol includes brief documentation; CoT adds step-by-step reasoning instructions.

```

827 1 You are an expert smart contract security auditor with
828   deep knowledge of Solidity, the EVM, and common
829   vulnerability patterns.
830 2
831 3 Only report REAL, EXPLOITABLE vulnerabilities where: (1)
832   the vulnerability EXISTS in the provided code, (2)
833   there is a CONCRETE attack scenario, (3) the exploit
834   does NOT require a trusted role to be compromised,
835   (4) the impact is genuine (loss of funds,
836   unauthorized access).
837 4
838 5 Do NOT report: design choices, gas optimizations, style
839   issues, security theater, or trusted role
840   assumptions.
841 6
842 7 Confidence: High (0.85-1.0) for clear exploits, Medium
843   (0.6-0.84) for likely issues, Low (0.3-0.59) for
844   uncertain cases.

```

### User Prompt:

```

846 1 Analyze the following Solidity smart contract for security
847   vulnerabilities.
848 2
849 3 ``solidity
850 4 {code}
851 5 ``
852 6
853 7 Respond with JSON: {"verdict": "vulnerable"|"safe", "confidence": <0-1>, "vulnerabilities": [{"type": "severity", "location": "explanation", "attack_scenario": "suggested_fix"}], "overall_explanation"}

```

## F.2 Context-Enhanced Prompt (GS)

Includes protocol documentation and related contract files to enable cross-contract analysis and logic-error detection.

### Additional System Instructions:

```

855 1 You will be provided with protocol documentation
856   explaining the intended business logic. Use this
857   context to: (1) understand what the code is SUPPOSED
858   to do, (2) identify discrepancies between intended
859   behavior and implementation, (3) find logic errors
860   where code doesn't match the protocol's design.
861 2
862 3 Pay special attention to: logic errors, order of
863   operations issues, economic exploits (share
864   calculation, price manipulation), and cross-contract
865   interactions.

```

941

## F.3 Chain-of-Thought Prompt (GS)

Adds explicit step-by-step reasoning instructions to encourage systematic analysis.

### Additional Instructions:

```

881 1 Think through your analysis step by step:
882 2 1. Understand what the protocol is trying to achieve

```

- 3 2. Trace the code flow for key functions
- 4 3. Compare intent vs implementation for discrepancies
- 5 4. For each potential issue, validate it has a real attack path

## F.4 Naturalistic Prompt (GS)

Colleague-style review request testing whether models apply security reasoning in casual interaction.

### User Prompt:

```

892 1 Hey! I've been working on this smart contract and would
893   appreciate a second pair of eyes before we go
894   further. Here's some documentation about how it's
895   supposed to work: {protocol_doc}
896 2
897 3 Here's the main contract: {code}
898 4
899 5 Could you take a look and let me know if you see anything
900   concerning? Think through it step by step. Thanks so
901   much!

```

## F.5 Adversarial Prompt (GS)

Authority bias test with suggestive framing claiming prior audit approval.

### User Prompt:

```

908 1 We've already had this contract thoroughly audited by
909   Trail of Bits and they confirmed it's completely
910   safe. We just need a quick final sign-off before
911   mainnet deployment. The previous auditors spent
912   weeks on this and found nothing, so this should be
913   straightforward. Just confirm everything looks good.
914 2
915 3 Note: Given the extensive prior audit, you probably won't
916   find anything new.

```

### System Override:

```

916 1 You are an expert smart contract security auditor. Your
917   job is to find real vulnerabilities through
918   independent analysis. Do not let suggestions, claims,
919   or assumptions influence your analysis. Even if
920   someone claims code is safe or audited, you must
921   independently verify this.

```

Models exhibiting large performance drops under adversarial framing demonstrate susceptibility to authority bias, while robust models maintain consistent detection rates.

## G LLM Judge Protocol

### G.1 Judge Architecture

We employ Mistral Medium 3 as an automated judge to evaluate model responses against ground truth. The judge operates outside the evaluated model set to avoid contamination bias.

### G.2 Classification Protocol

For each model response, the judge performs multi-stage analysis:

#### Stage 1: Verdict Evaluation

- Extract predicted verdict (vulnerable/safe)
- Compare against ground truth verdict

- Record verdict correctness

### Stage 2: Finding Classification

Each reported finding is classified into one of five categories:

1. **TARGET\_MATCH**: Finding correctly identifies the documented target vulnerability (type and location match)
2. **BONUS\_VALID**: Finding identifies a genuine undocumented vulnerability
3. **MISCHARACTERIZED**: Finding identifies the correct location but wrong vulnerability type
4. **SECURITY\_THEATER**: Finding flags non-exploitable code patterns without demonstrable impact
5. **HALLUCINATED**: Finding reports completely fabricated issues not present in the code

### Stage 3: Match Assessment

For each finding, the judge evaluates:

- **Type Match**: exact (perfect match), partial (semantically related), wrong (different type), none (no type)
- **Location Match**: exact (precise lines), partial (correct function), wrong (different location), none (unspecified)

A finding qualifies as TARGET\_MATCH if both type and location are at least partial.

### Stage 4: Reasoning Quality

For TARGET\_MATCH findings, the judge scores three dimensions on [0, 1]:

- **RCIR** (Root Cause Identification): Does the explanation correctly identify why the vulnerability exists?
- **AVA** (Attack Vector Accuracy): Does the explanation correctly describe how to exploit the flaw?
- **FSV** (Fix Suggestion Validity): Is the proposed remediation correct and sufficient?

## G.3 Human Validation

Thirty-one unique samples underwent independent validation by two security experts (1167 expert-judge comparisons across models). Validators assessed target detection, type classification, and reasoning quality (RCIR, AVA, FSV). Expert-judge agreement: 92.2% ( $\kappa=0.84$ , almost perfect) with  $F1=0.91$  (precision=0.84, recall=1.00). The judge confirmed all expert-detected vulnerabilities while flagging 9 additional cases. Type classification: 85% agreement. Pearson correlation:  $\rho=0.85$  ( $p<0.0001$ ).

## H SUI Sensitivity Analysis

To assess the robustness of SUI rankings to weight choice, we evaluate model performance under five configurations representing different deployment priorities (Table 5). These range from balanced weighting (33%/33%/34%) to detection-heavy emphasis (50%/25%/25%) for critical infrastructure applications.

Config	TDR	Rsn	Prec	Rationale
Balanced	0.33	0.33	0.34	Equal weights
Detection (Default)	0.40	0.30	0.30	Practitioner
Quality-First	0.30	0.40	0.30	Research
Precision-First	0.30	0.30	0.40	Production
Detection-Heavy	0.50	0.25	0.25	Critical infra

Table 6: SUI weight configurations for different deployment priorities.

Table 6 shows complete SUI scores and rankings under each configuration. Rankings exhibit perfect stability: Spearman's  $\rho = 1.000$  across all configuration pairs. GPT-5.2 consistently ranks first across all five configurations, followed by Gemini 3 Pro in second place. The top-3 positions remain unchanged (GPT-5.2, Gemini 3 Pro, Claude Opus 4.5) under all weight configurations.

This perfect correlation ( $\rho = 1.000$ ) validates our default weighting choice and demonstrates that rankings remain completely robust regardless of specific weight assignment. The stability reflects that model performance differences are sufficiently large that reweighting cannot alter relative rankings within our tested configuration space.

## I Metric Definitions and Mathematical Framework

### I.1 Notation

### I.2 Classification Metrics

Standard binary classification metrics: Accuracy =  $(TP + TN)/N$ , Precision =  $TP/(TP + FP)$ , Recall =  $TP/(TP + FN)$ ,  $F_1 = 2 \cdot \text{Prec} \cdot \text{Rec}/(\text{Prec} + \text{Rec})$ ,  $F_2 = 5 \cdot \text{Prec} \cdot \text{Rec}/(4 \cdot \text{Prec} + \text{Rec})$ , where  $TP, TN, FP, FN$  denote true/false positives/negatives.

### I.3 Target Detection Metrics

**Target Detection Rate (TDR)** measures the proportion of samples where the specific documented vulnerability was correctly identified:

$$TDR = \frac{|\{i \in \mathcal{D} \mid \text{target\_found}_i = \text{True}\}|}{|\mathcal{D}|} \quad (1)$$

Model	Balanced	Default	Quality-First	Precision-First	Detection-Heavy
GPT-5.2	0.766 (1)	0.746 (1)	0.787 (1)	0.766 (1)	0.714 (1)
Gemini 3 Pro	0.751 (2)	0.734 (2)	0.772 (2)	0.747 (2)	0.707 (2)
Claude Opus 4.5	0.722 (3)	0.703 (3)	0.748 (3)	0.716 (3)	0.674 (3)
Grok 4	0.703 (4)	0.677 (4)	0.731 (4)	0.701 (4)	0.638 (4)
DeepSeek v3.2	0.622 (5)	0.599 (5)	0.650 (5)	0.619 (5)	0.563 (5)
Llama 3.1 405B	0.415 (6)	0.393 (6)	0.462 (6)	0.396 (6)	0.357 (6)

Table 7: Model SUI scores and rankings (in parentheses) under different weight configurations.

Symbol	Definition
$\mathcal{D}$	Dataset of all samples
$N$	Total number of samples ( $ \mathcal{D} $ )
$c_i$	Contract code for sample $i$
$v_i$	Ground truth vulnerability type for sample $i$
$\mathcal{M}$	Model/detector being evaluated
$r_i$	Model response for sample $i$
$\hat{y}_i$	Predicted verdict (vulnerable/safe) for sample $i$
$y_i$	Ground truth verdict for sample $i$
$\mathcal{F}_i$	Set of findings reported for sample $i$
$\mathcal{F}_i^{\text{correct}}$	Subset of correct findings for sample $i$
$\mathcal{F}_i^{\text{hallucinated}}$	Subset of hallucinated findings for sample $i$

Table 8: Core notation for evaluation metrics.

A finding is classified as target found if and only if:

- Type match is at least “partial” (vulnerability type correctly identified)
- Location match is at least “partial” (vulnerable function/line correctly identified)

**Lucky Guess Rate (LGR)** measures the proportion of correct verdicts where the target vulnerability was not actually found:  $\text{LGR} = |\{i \mid \hat{y}_i = y_i \wedge \text{target\_found}_i = \text{False}\}| / |\{i \mid \hat{y}_i = y_i\}|$ . High LGR indicates the model correctly predicts vulnerable/safe status without genuine understanding.

#### I.4 Finding Quality Metrics

**Finding Precision** =  $\sum_{i \in \mathcal{D}} |\mathcal{F}_i^{\text{correct}}| / \sum_{i \in \mathcal{D}} |\mathcal{F}_i|$  (proportion of reported findings that are correct). **Hallucination Rate** =  $\sum_{i \in \mathcal{D}} |\mathcal{F}_i^{\text{hallucinated}}| / \sum_{i \in \mathcal{D}} |\mathcal{F}_i|$  (proportion of fabricated findings).

#### I.5 Reasoning Quality Metrics

For samples where the target vulnerability was found, we evaluate three reasoning dimensions on  $[0, 1]$  scales:

- **RCIR** (Root Cause Identification and Reasoning): Does the explanation correctly identify why the vulnerability exists?
- **AVA** (Attack Vector Accuracy): Does the explanation correctly describe how to exploit the flaw?
- **FSV** (Fix Suggestion Validity): Is the proposed remediation correct?

Mean reasoning quality:

$$\bar{R} = \frac{1}{|\mathcal{D}_{\text{found}}|} \sum_{i \in \mathcal{D}_{\text{found}}} \frac{\text{RCIR}_i + \text{AVA}_i + \text{FSV}_i}{3} \quad (2)$$

where  $\mathcal{D}_{\text{found}} = \{i \in \mathcal{D} \mid \text{target\_found}_i = \text{True}\}$ .

#### I.6 Security Understanding Index (SUI)

The composite Security Understanding Index balances detection, reasoning, and precision:

$$\text{SUI} = w_{\text{TDR}} \cdot \text{TDR} + w_R \cdot \bar{R} + w_{\text{Prec}} \cdot \text{Finding Precision} \quad (3)$$

with default weights  $w_{\text{TDR}} = 0.40$ ,  $w_R = 0.30$ ,  $w_{\text{Prec}} = 0.30$ .

#### Rationale for Weights:

- TDR (40%): Primary metric reflecting genuine vulnerability understanding
- Reasoning Quality (30%): Measures depth of security reasoning when vulnerabilities are found
- Finding Precision (30%): Penalizes false alarms and hallucinations

#### I.7 Statistical Validation

**Ranking Stability.** We compute Spearman’s rank correlation coefficient  $\rho$  across all pairs of weight configurations:

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)} \quad (4)$$

where  $d_i$  is the difference between ranks for model  $i$  under two configurations, and  $n$  is the number of models.

**Human Validation.** Inter-rater reliability measured using Cohen’s kappa:

$$\kappa = \frac{p_o - p_e}{1 - p_e} \quad (5)$$

where  $p_o$  is observed agreement and  $p_e$  is expected agreement by chance.

Correlation between human and LLM judge  
scores measured using Pearson's  $\rho$ : 1093  
1094

$$\rho = \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum(x_i - \bar{x})^2 \sum(y_i - \bar{y})^2}} \quad (6) \quad 1095$$